

Constraint Satisfaction Problems (CSPs)

- In Uninformed/Informed Search the objective is:
 - Planning: sequences of actions
 - The path to the goal is the important thing
 - Paths have various costs, depths
 - Heuristics give problem-specific guidance
- CSP is a different (special) kind of search problem:
 - Identification: assignments to variables
 - The goal itself is important, not the path
 - All paths at the same depth (for some formulations)
 - CSPs are specialized for identification problems

Identification in essence is
“Finding/Discovering the Goal”

- Standard search problems: (Uniformed/Informed Search)
 - State is a “black box”: arbitrary data structure
 - Goal test can be any function over states
 - Successor function can also be anything
- Constraint satisfaction problems (CSPs):
 - A special subset of search problems
 - State is defined by variables X_i with values from a domain D (sometimes D depends on i)
 - Goal test is a set of constraints specifying allowable combinations of values for subsets of variables
- Simple example of a *formal representation language*
- Allows useful general-purpose algorithms with more power than standard search algorithms

Example: Map Coloring

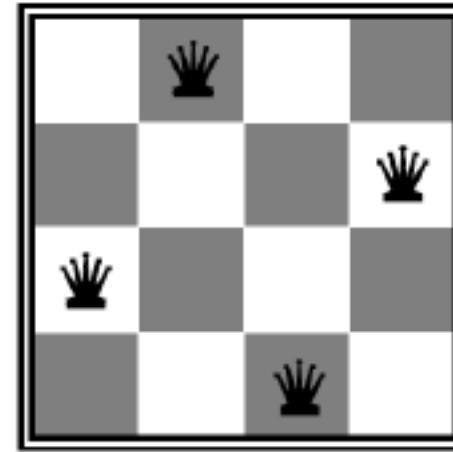
- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D = \{\text{red, green, blue}\}$
- Constraints: adjacent regions must have different colors
 - Implicit: $\text{WA} \neq \text{NT}$
 - Explicit: $(\text{WA}, \text{NT}) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$



- Solutions are assignments satisfying all constraints, e.g.:
 $\{\text{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green}\}$

Example: N-Queens

- Formulation 1:
 - Variables: X_{ij}
 - Domains: $\{0, 1\}$
 - Constraints



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j+k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k, j-k}) \in \{(0, 0), (0, 1), (1, 0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- Formulation 2:

- Variables: Q_k

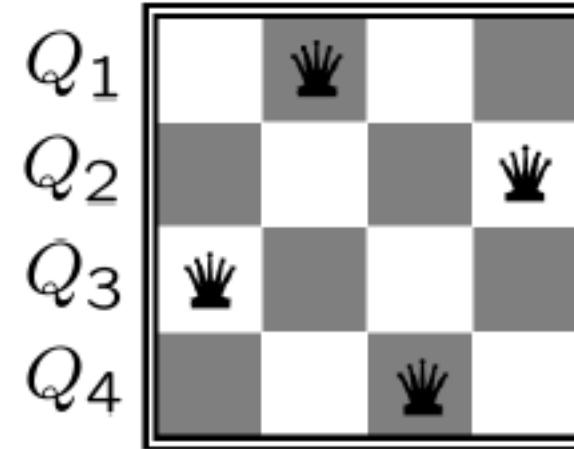
- Domains: $\{1, 2, 3, \dots, N\}$

- Constraints:

Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

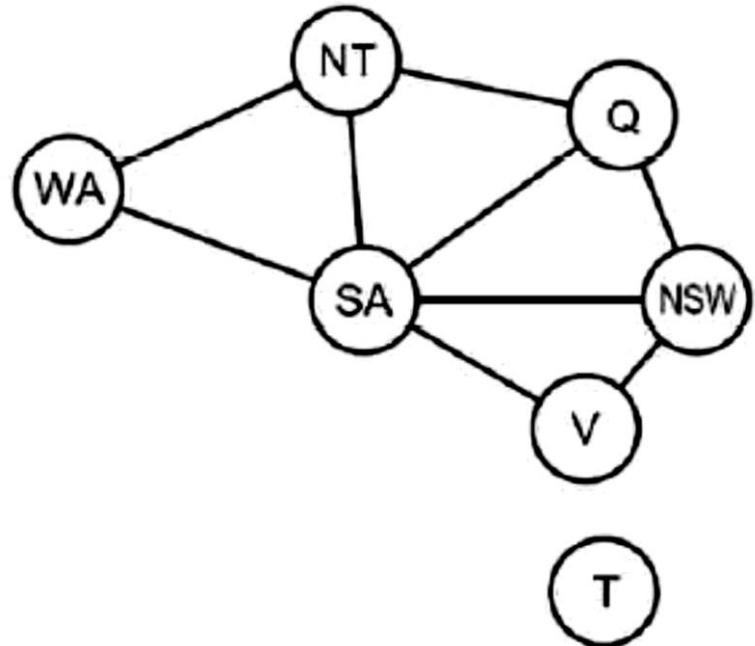
Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

• • •



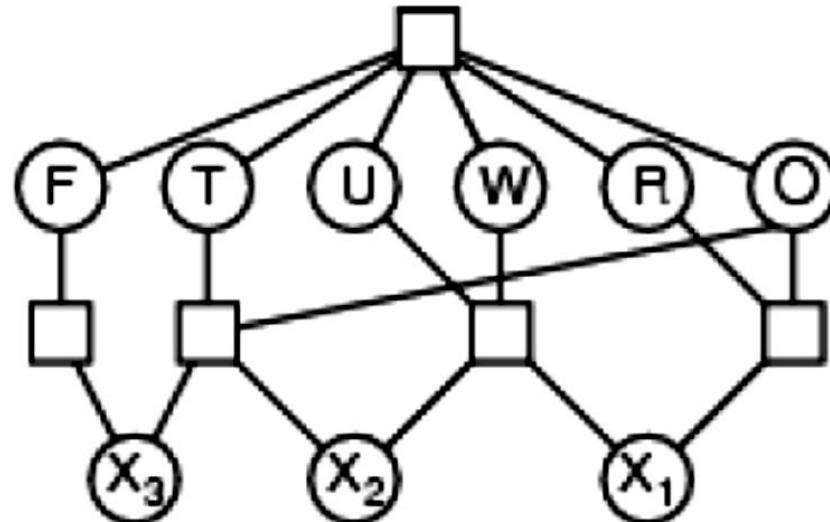
Constraint Graphs

- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints
- General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!



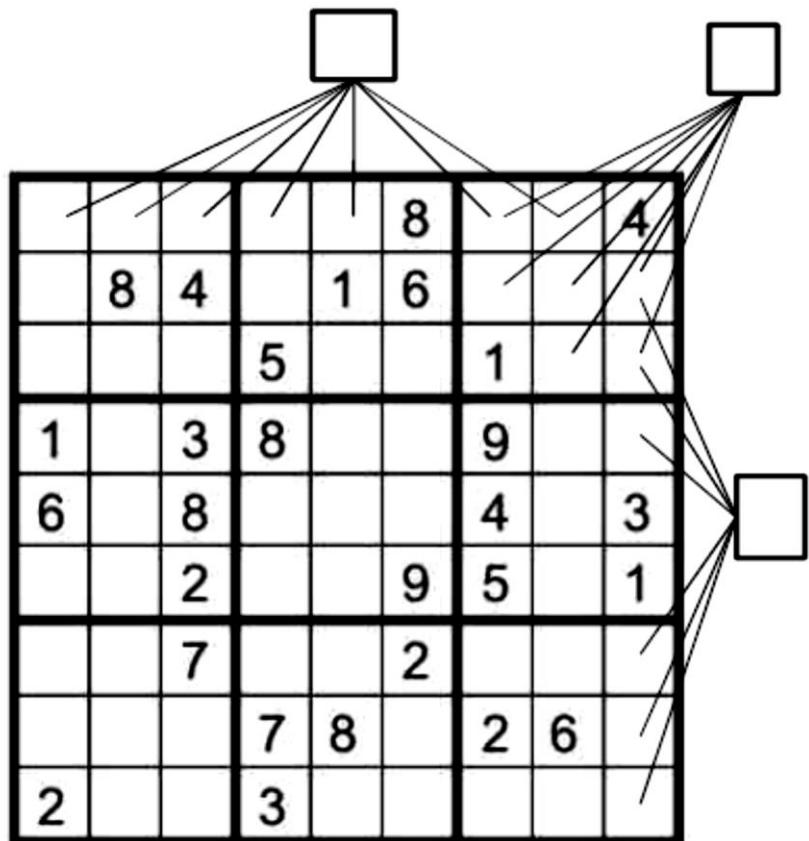
Cryptarithmetic

$$\begin{array}{r} \text{T} \ \text{W} \ \text{O} \\ + \ \text{T} \ \text{W} \ \text{O} \\ \hline \text{F} \ \text{O} \ \text{U} \ \text{R} \end{array}$$



- Variables: $F, T, U, W, R, O, X_1, X_2, X_3$
- Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Constraints: $\text{Alldiff}(F, T, U, W, R, O)$
 - $O + O = R + 10 \cdot X_1$
 - $X_1 + W + W = U + 10 \cdot X_2$
 - $X_2 + T + T = O + 10 \cdot X_3$
 - $X_3 = F, T \neq 0, F \neq 0$

Sudoku



- Variables:
 - Each (open) square
- Domains:
 - $\{1, 2, \dots, 9\}$
- Constraints:
 - 9-way alldiff for each column
 - 9-way alldiff for each row
 - 9-way alldiff for each region
 - (or can have a bunch of pairwise inequality constraints)

Solution: a value for each cell satisfying the constraints

Scheduling

- Want to schedule a time and a space for each final exam so that no student is scheduled to take more than one final at the same time.
- The space allocated has to be available at the time set.
- The space has to be large enough to accommodate all of the students taking the exam.

Varieties of CSPs

- Discrete Variables
 - Finite domains
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
 - Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Linear constraints solvable, nonlinear undecidable
- Continuous variables
 - Linear constraints solvable in polynomial time by LP methods

Varieties of Constraints

- **Varieties of Constraints**
 - Unary constraints involve a single variable (equivalent to reducing domains), e.g.:
 $SA \neq \text{green}$
 - Binary constraints involve pairs of variables, e.g.:
 $SA \neq WA$
 - Higher-order constraints involve 3 or more variables:
e.g., cryptarithmetic column constraints
- **Preferences (soft constraints):**
 - E.g., red is better than green
 - Often representable by a cost for each variable assignment
 - Gives constrained optimization problems

Real-World CSPs

- Assignment problems: e.g., who teaches what class
- Timetabling problems: e.g., which class is offered when and where?
- Hardware configuration
- Transportation scheduling
- Factory scheduling
- Circuit layout
- Fault diagnosis
- ... lots more!

- Many real-world problems involve real-valued variables...

Solving CSPs

- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment, {}
 - Successor function: assign a value to an unassigned variable
 - Goal test: the current assignment is complete and satisfies all constraints

Solving CSPs (continued)

- A *state* is an *assignment* of values to some or all variables.
 - An assignment is *complete* when every variable has a value.
 - An assignment is *partial* when some variables have no values.
- ***Consistent assignment***
 - assignment does not violate the constraints
- A *solution* to a CSP is a complete and consistent assignment.
- Some CSPs require a solution that maximizes an *objective function*.

Backtracking Search

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
 - Variable assignments are commutative, so fix ordering
 - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
 - I.e. consider only values which do not conflict previous assignments
 - Might have to do some computation to check the constraints
 - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for $n \approx 25$

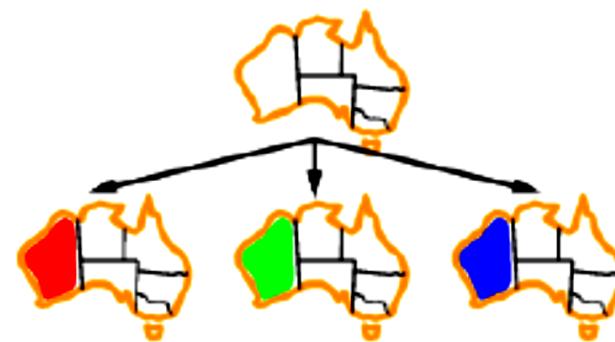
Backtracking Search – Pseudo-Code

```
function BACKTRACKING-SEARCH( csp) returns a solution, or failure
    return RECURSIVE-BACKTRACKING( {}, csp)
function RECURSIVE-BACKTRACKING( assignment, csp) returns a solution, or
failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE( Variables[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment according to Constraints[csp] then
            add { var = value } to assignment
            result  $\leftarrow$  RECURSIVE-BACKTRACKING( assignment, csp)
            if result  $\neq$  failure then return result
            remove { var = value } from assignment
    return failure
```

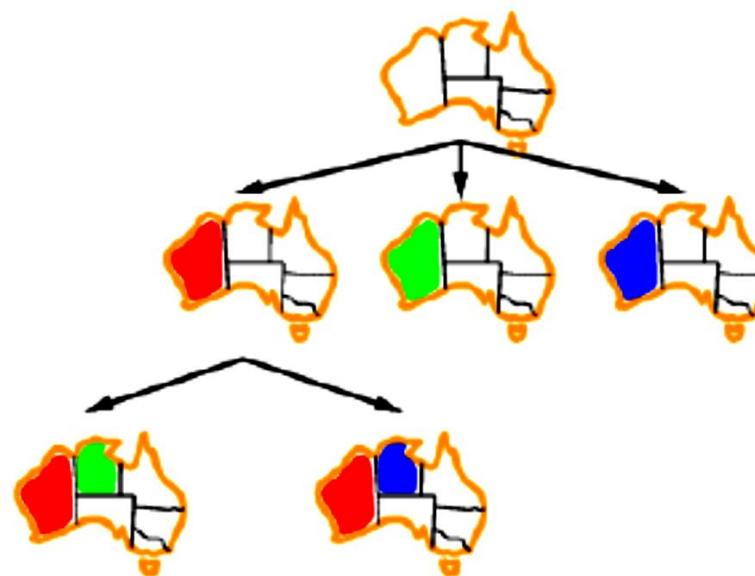
Backtracking example



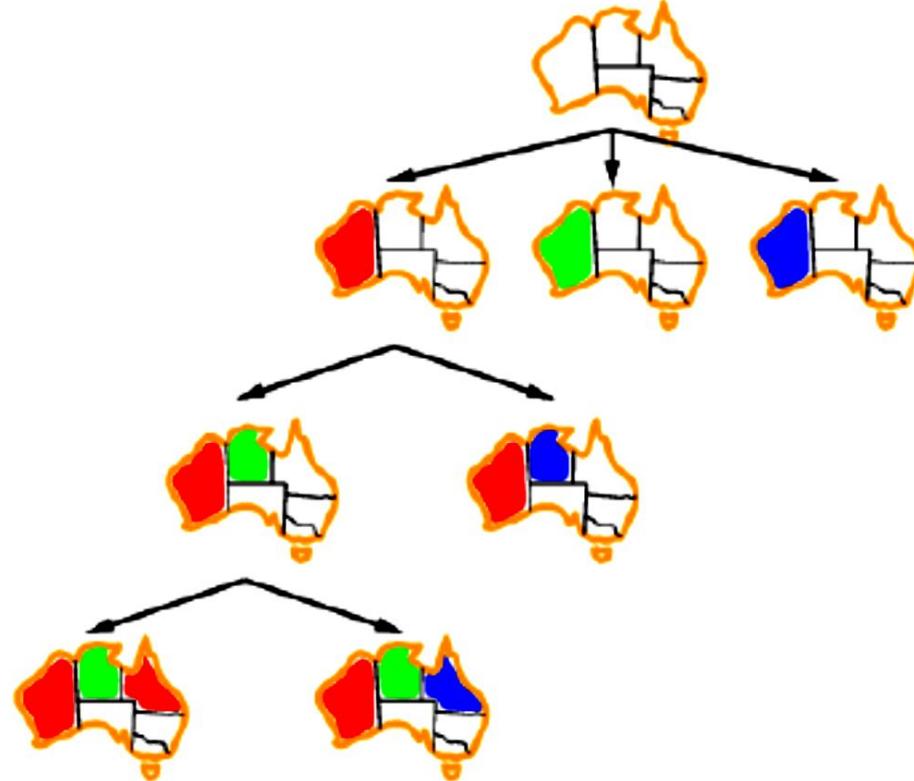
Backtracking example



Backtracking example



Backtracking example

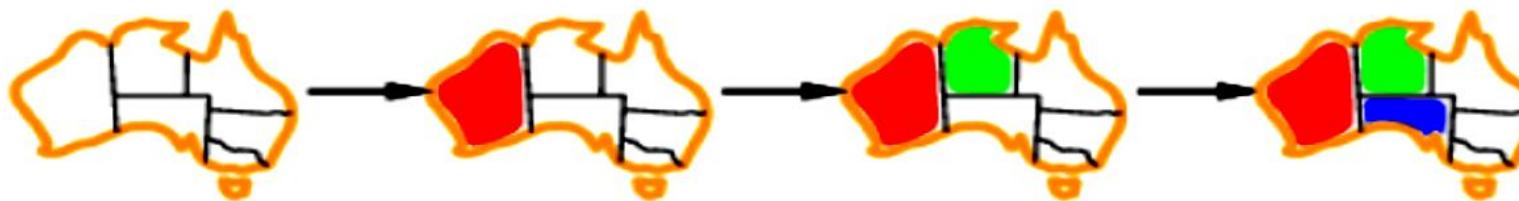


Improving Backtracking

- General-purpose ideas give huge gains in speed
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Filtering: Can we detect inevitable failure early?
- Structure: Can we exploit the problem structure?

Most constrained variable

- **Most constrained variable:**
choose the variable with the fewest legal values



- a.k.a. minimum remaining values (MRV) heuristic
- Fail-first heuristic
If a variable has no legal values that will be picked first and fail

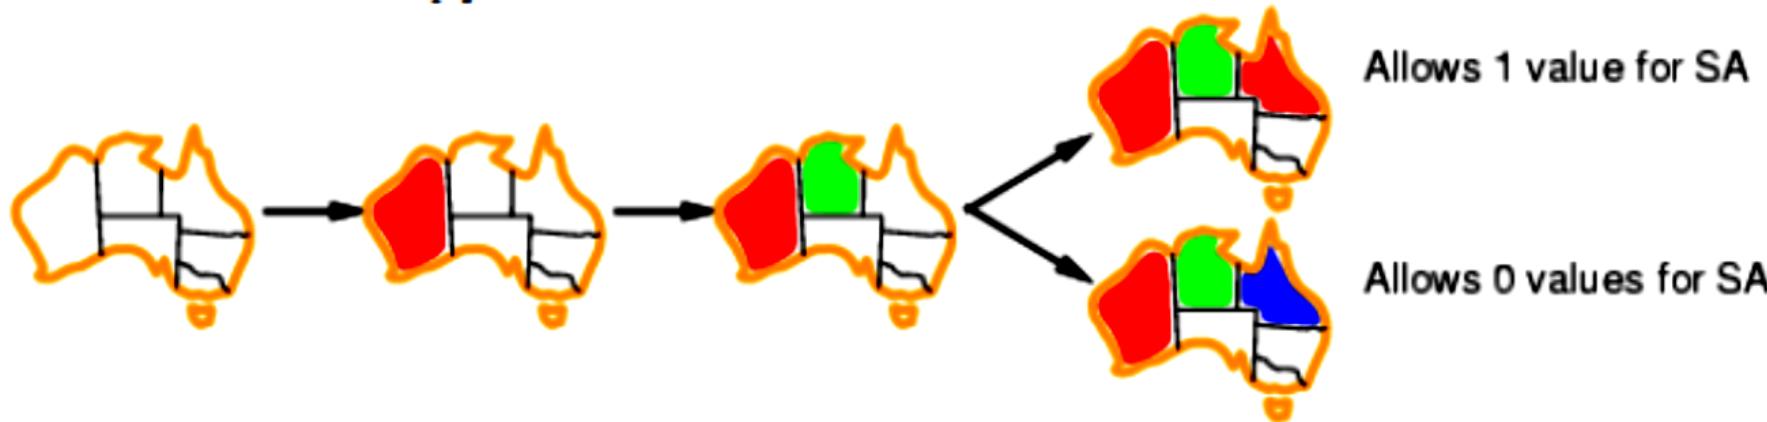
Most constraining variable

- Tie-breaker among most constrained variables
- Most constraining variable:
 - choose the variable with the most constraints on remaining variables



Least constraining value

- Given a variable, choose the least constraining value:
 - the one that rules out the fewest values in the remaining variables



- Combining these heuristics makes 1000 queens feasible

Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



WA

NT

q

NSW

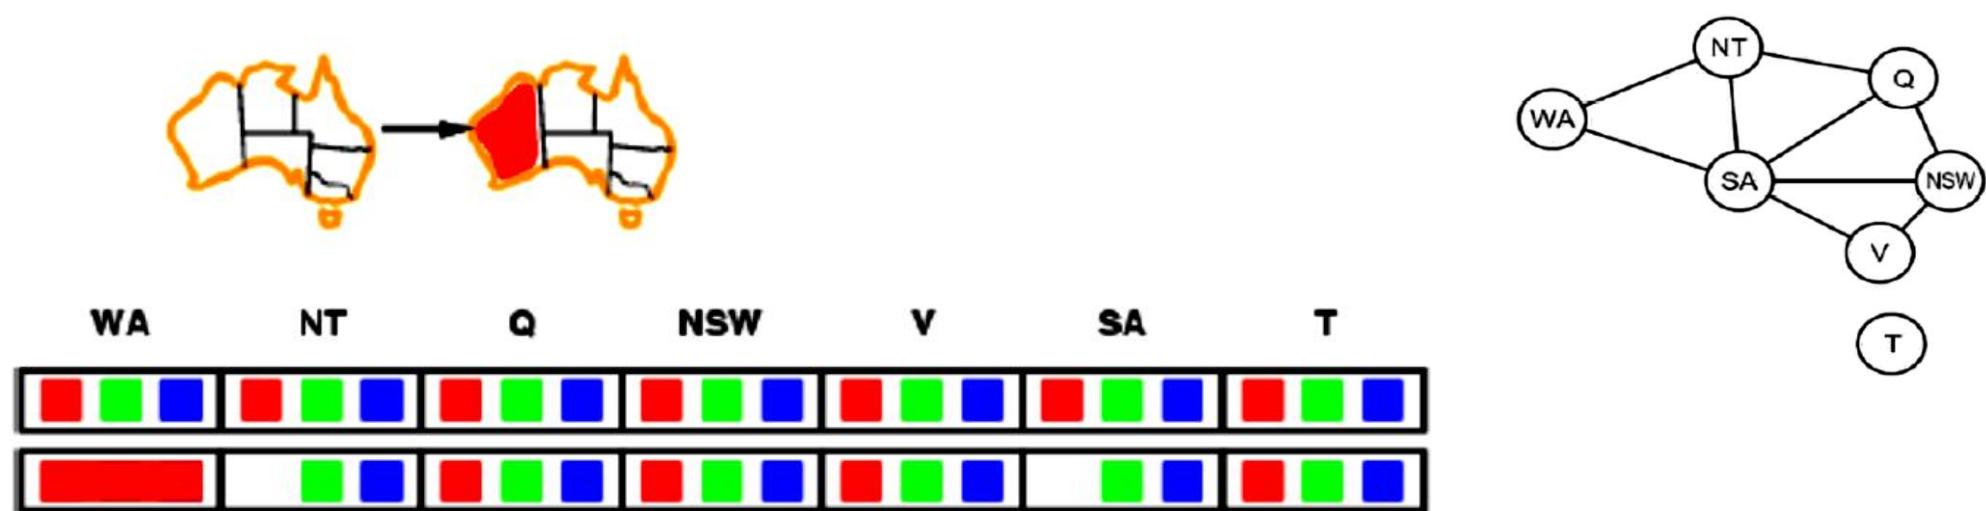
v

SA

T

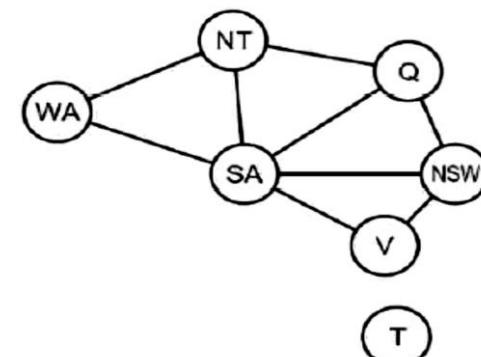
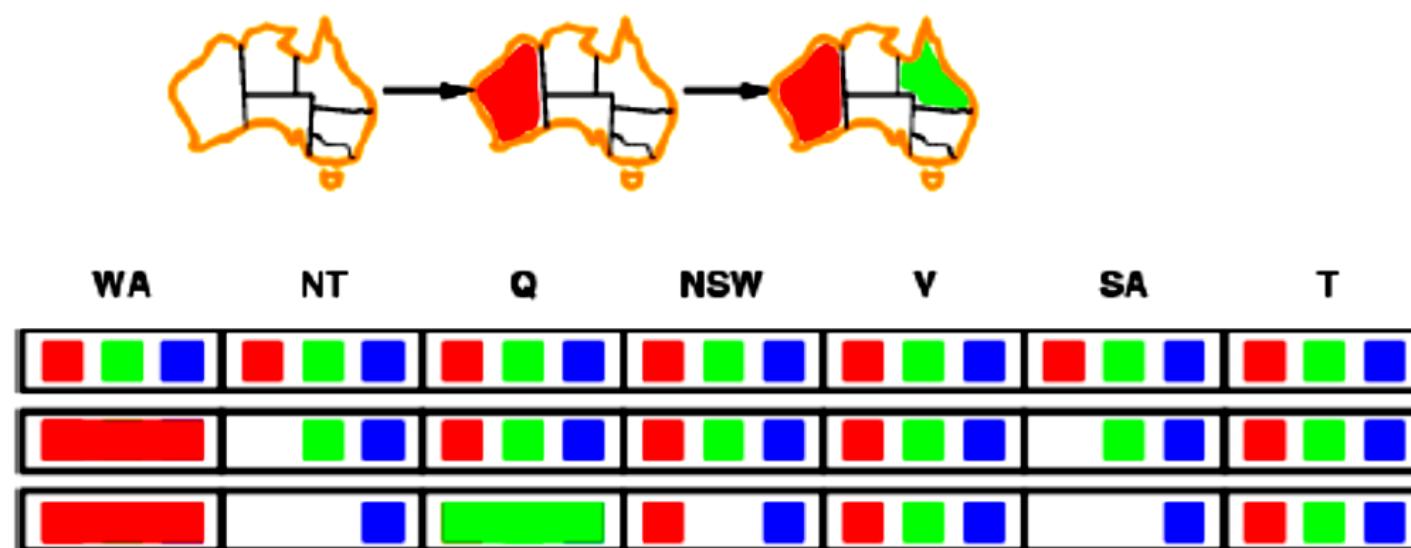
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



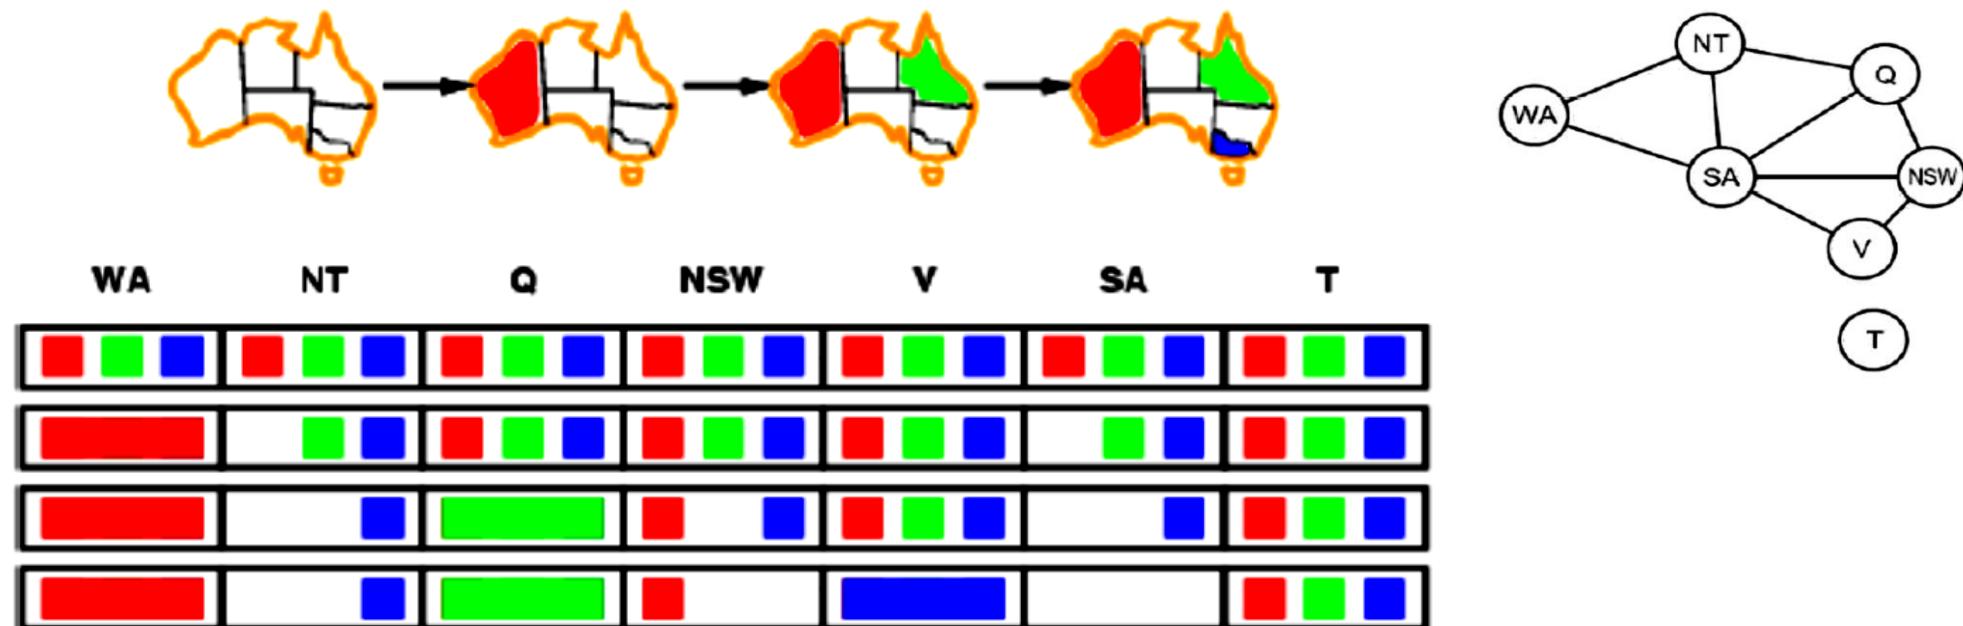
Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values

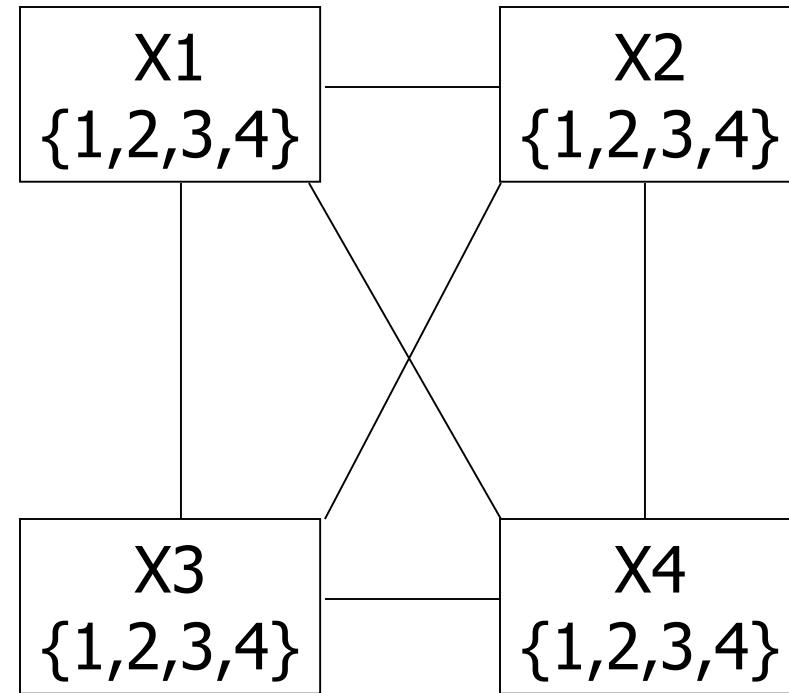
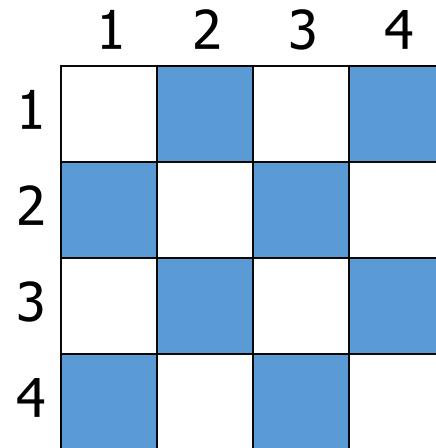


Forward checking

- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values

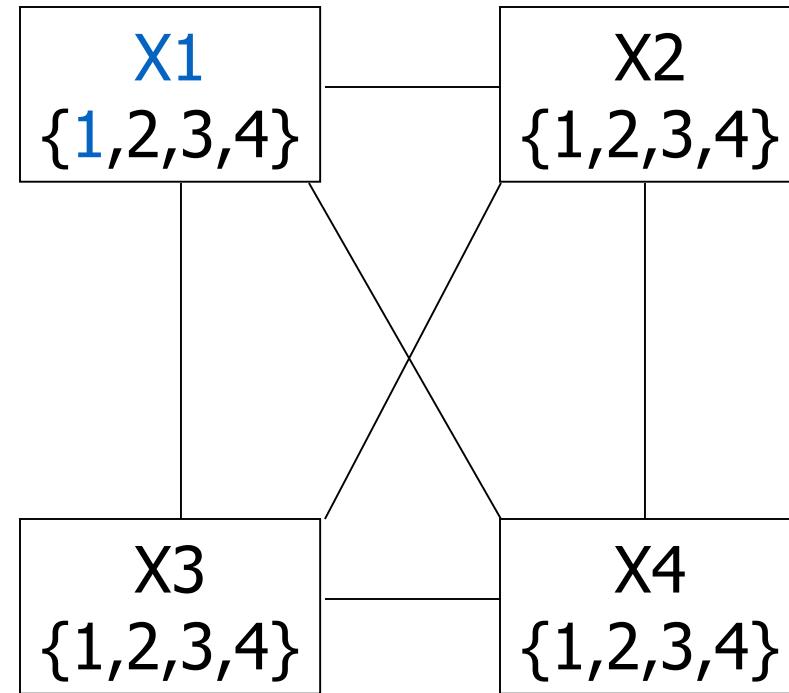
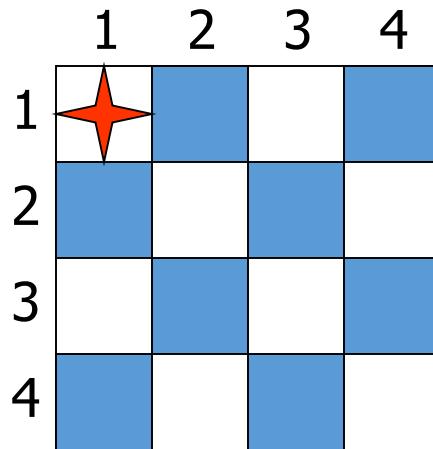


Another Forward Checking Example: 4-Queens Problem

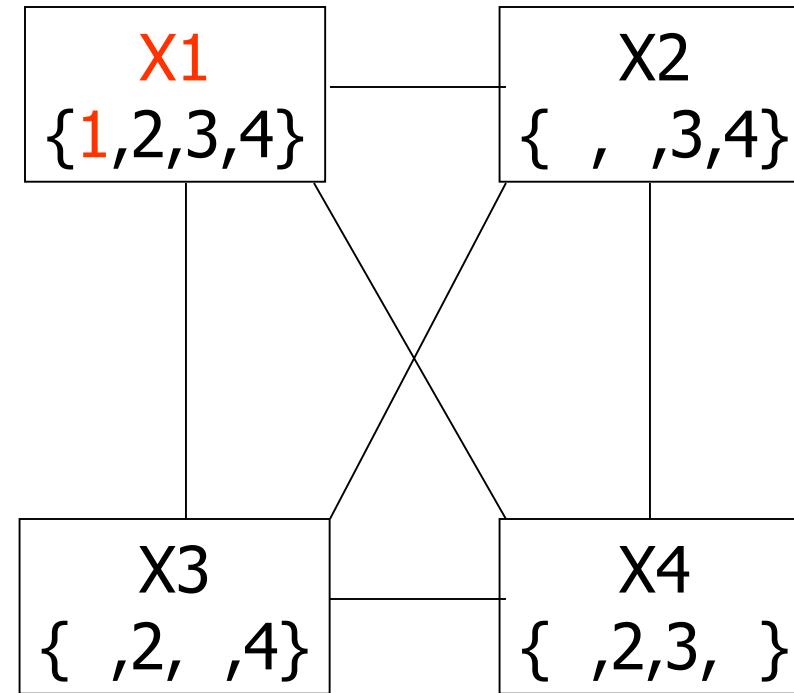
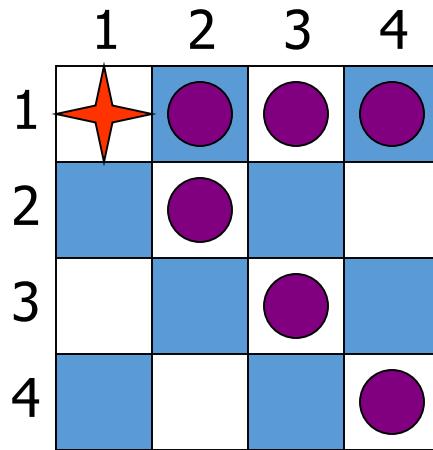


$X_i = i^{\text{th}}$ Queen can be placed anywhere in column i

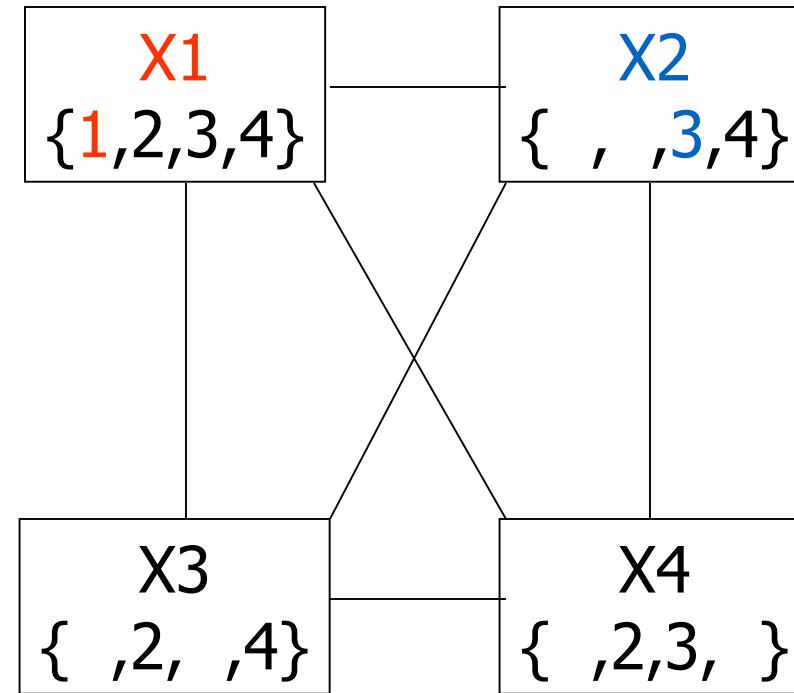
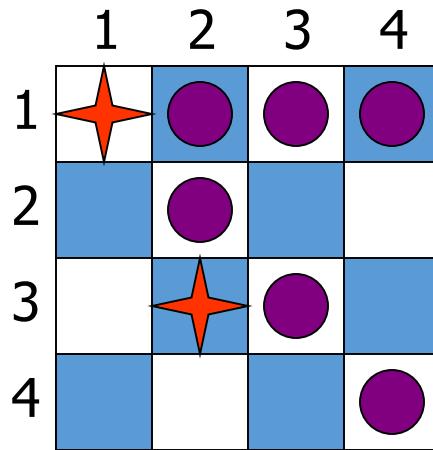
Example: 4-Queens Problem



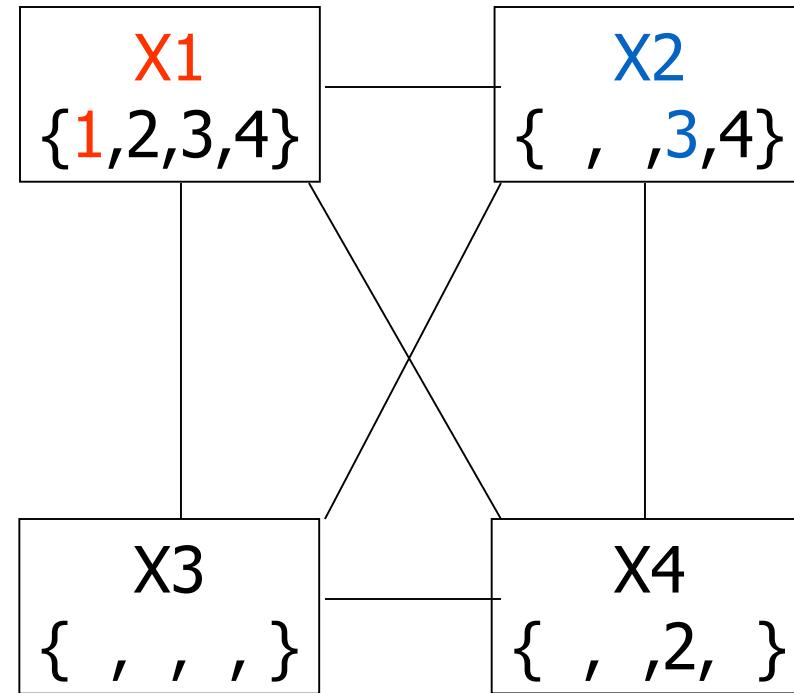
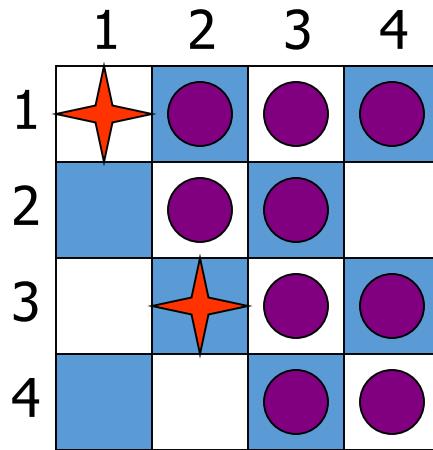
Example: 4-Queens Problem



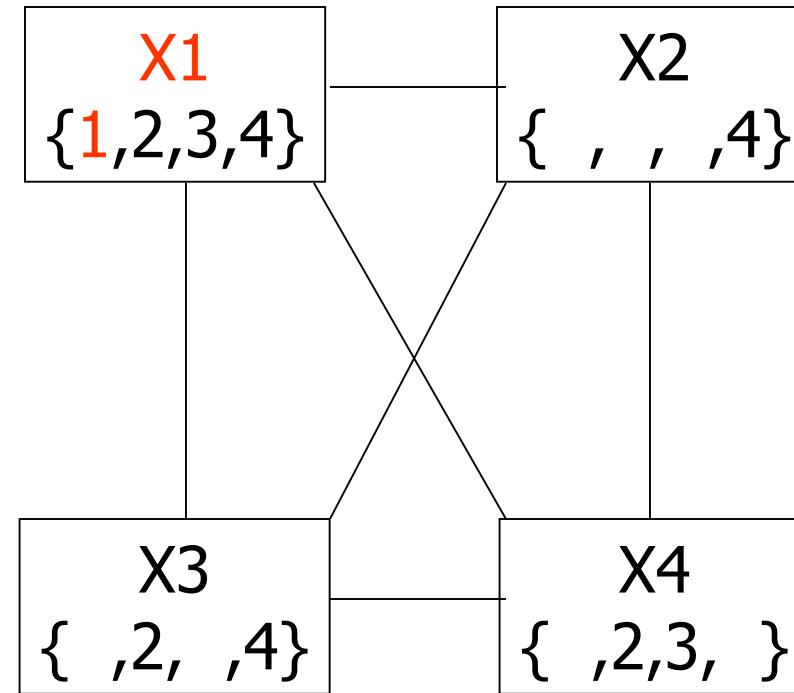
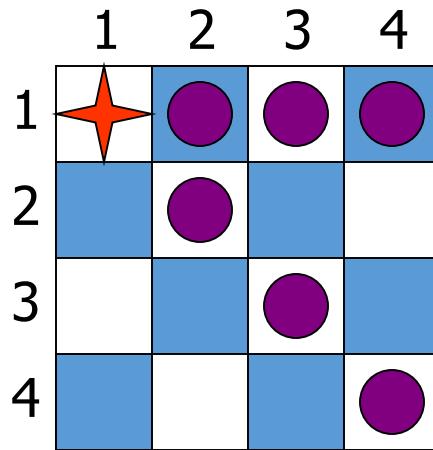
Example: 4-Queens Problem



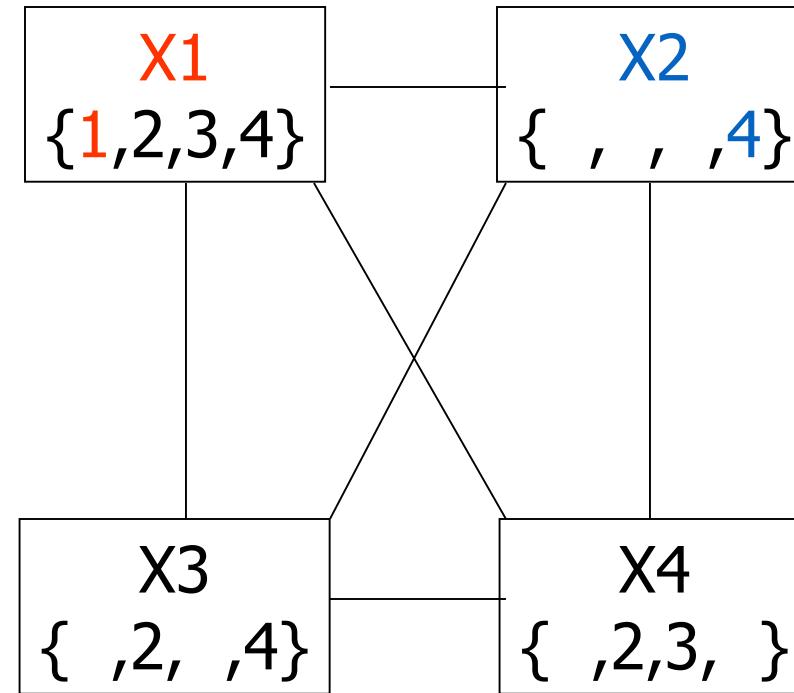
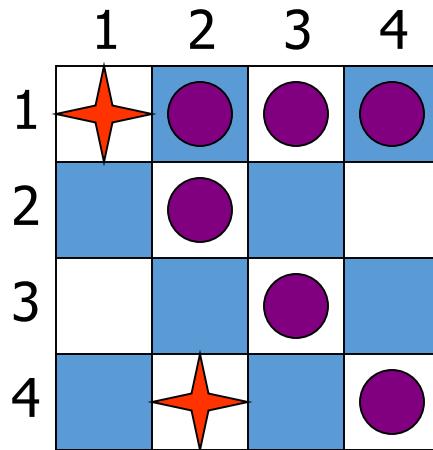
Example: 4-Queens Problem



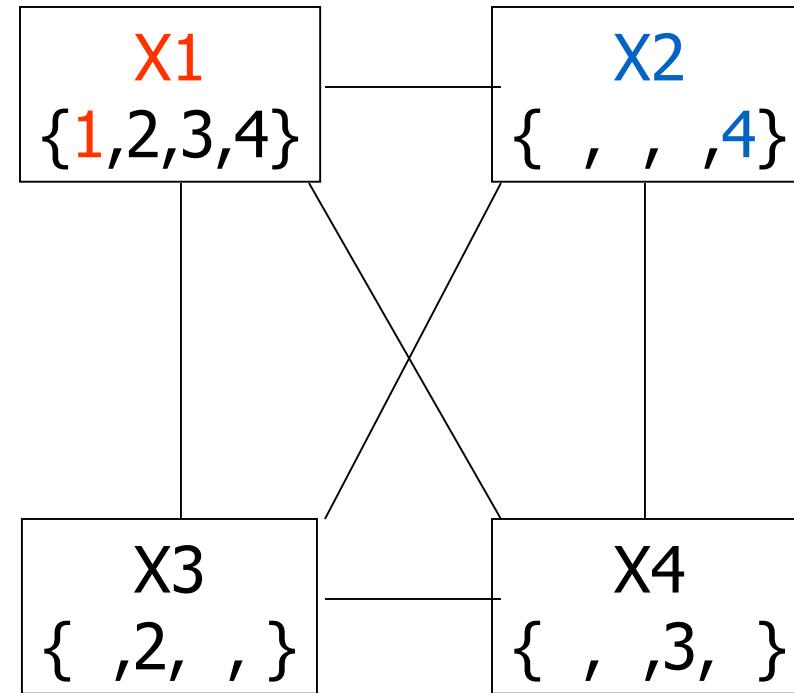
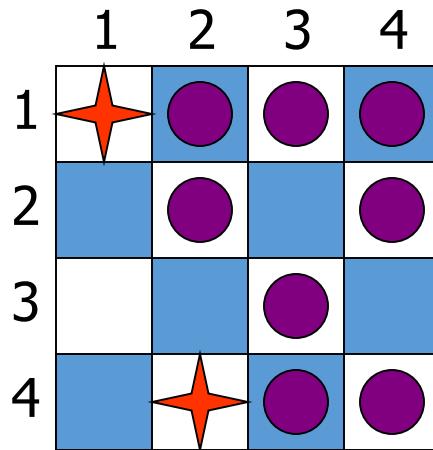
Example: 4-Queens Problem



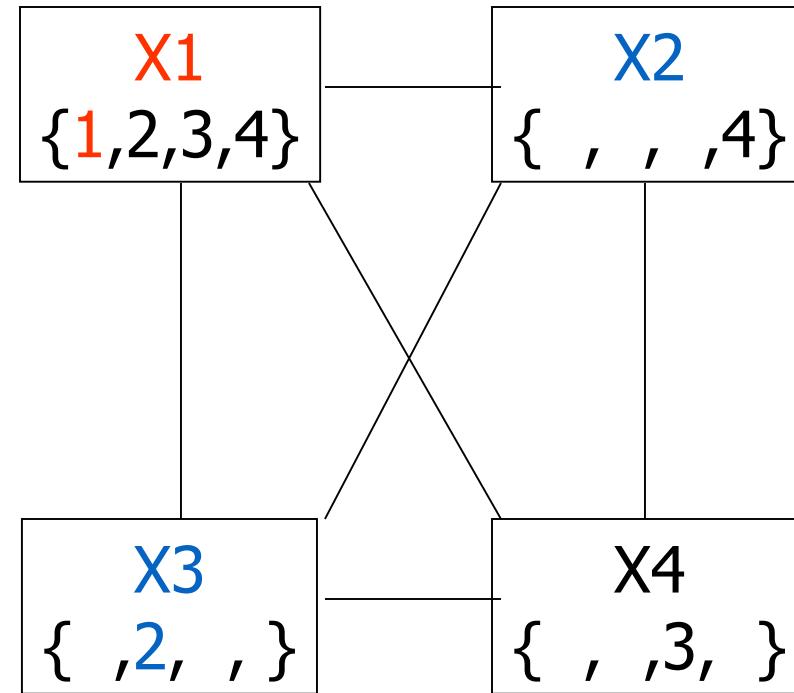
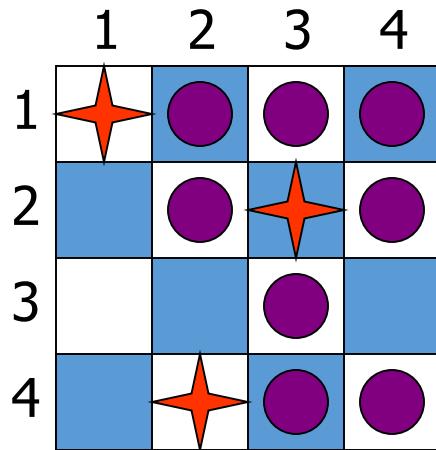
Example: 4-Queens Problem



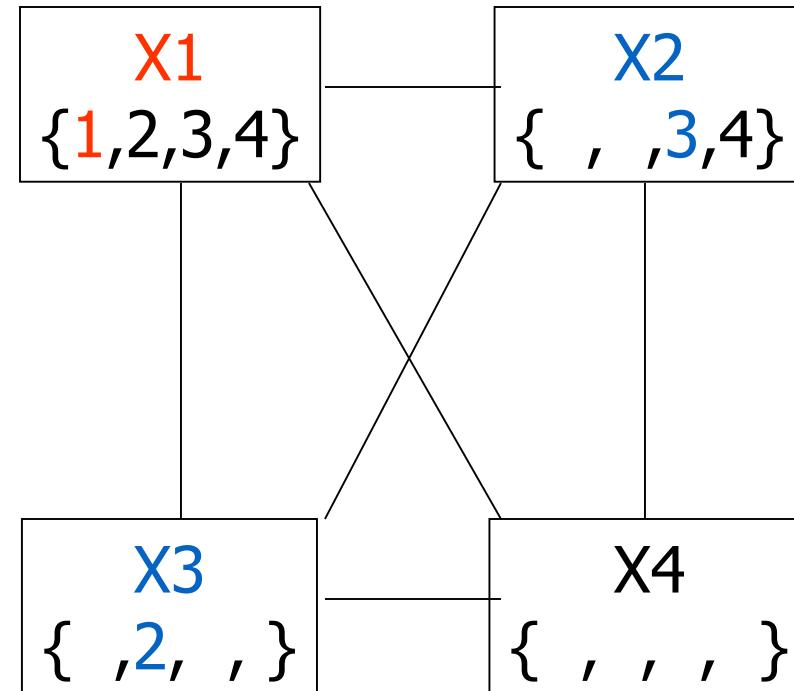
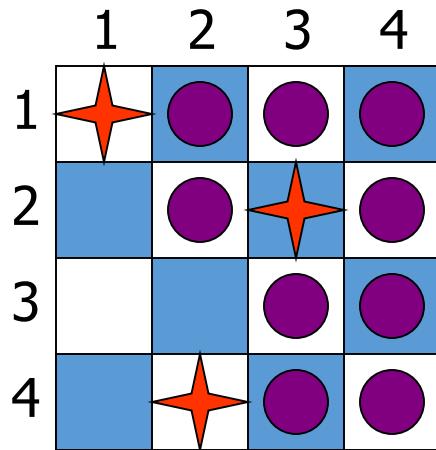
Example: 4-Queens Problem



Example: 4-Queens Problem



Example: 4-Queens Problem



Comparison of CSP algorithms on different problems

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV
USA	(> 1,000K)	(> 1,000K)	2K	60
n -Queens	(> 40,000K)	13,500K	(> 40,000K)	817K
Zebra	3,859K	1K	35K	0.5K
Random 1	415K	3K	26K	2K
Random 2	942K	27K	77K	15K

Median number of consistency checks over 5 runs to solve problem

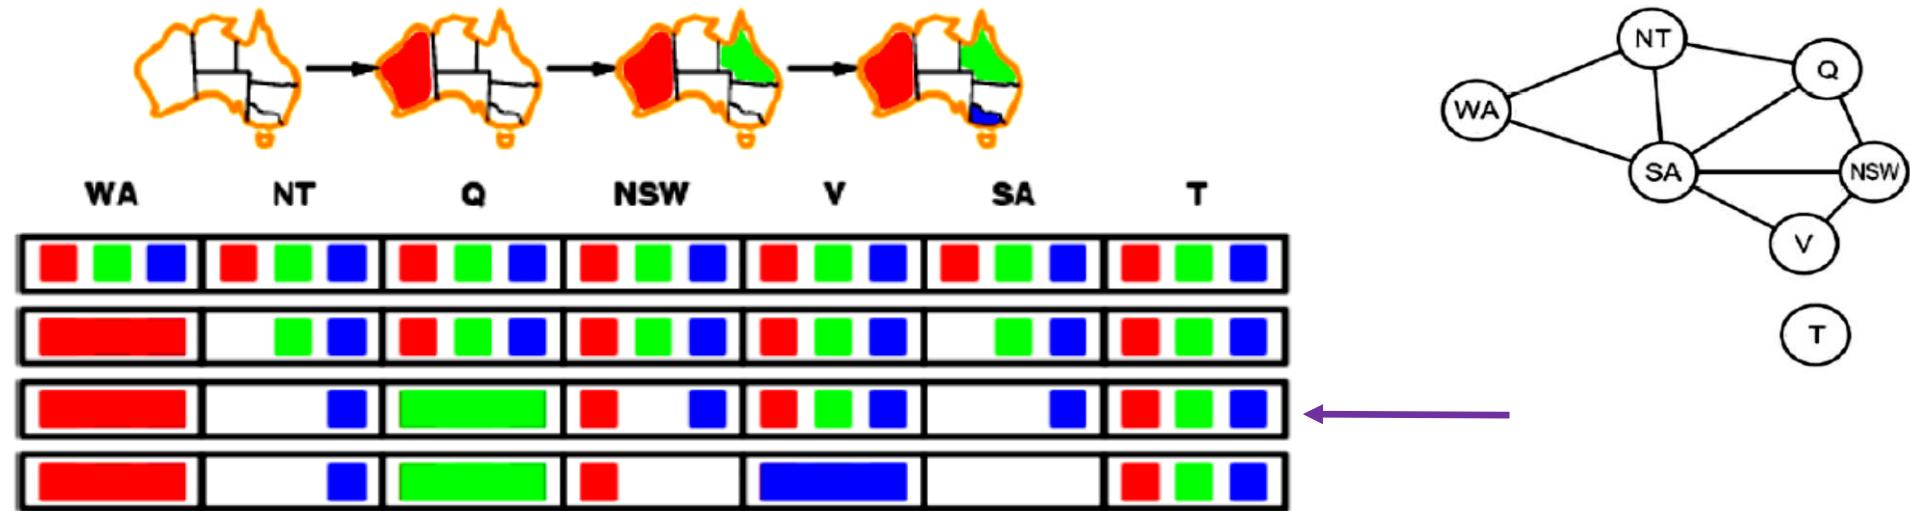
Parentheses -> no solution found

USA: 4 coloring

n -Queens: n ranges from 2 to 50

Constraint propagation

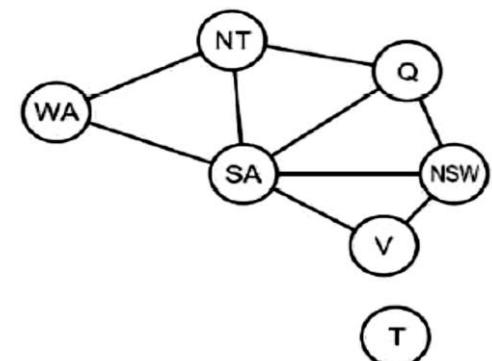
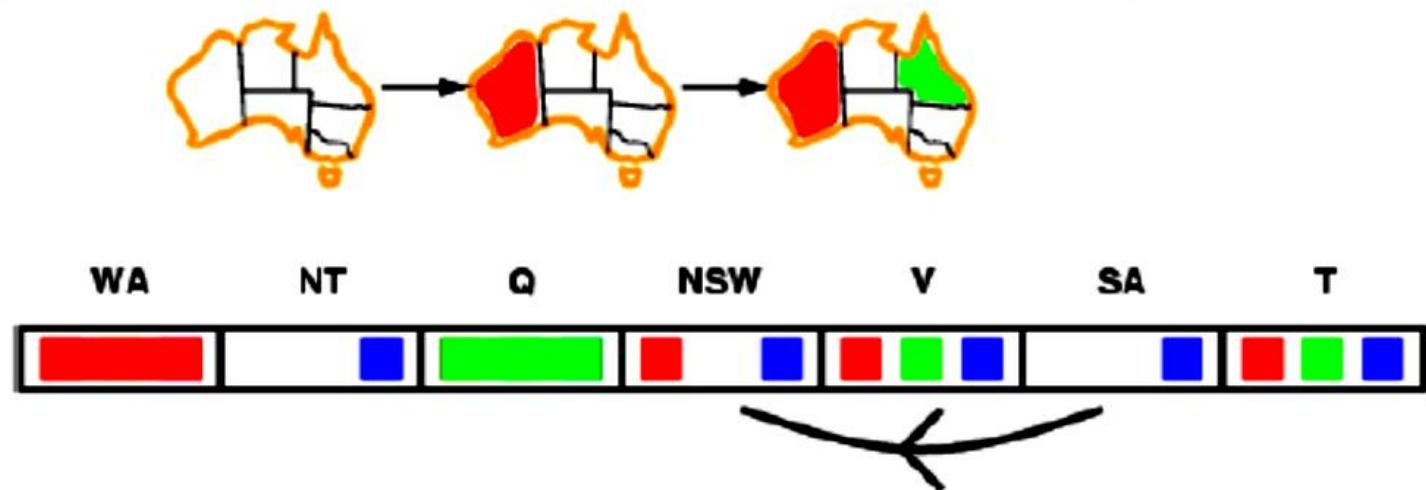
- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



- NT and SA cannot both be blue!
- Constraint propagation repeatedly enforces constraints locally

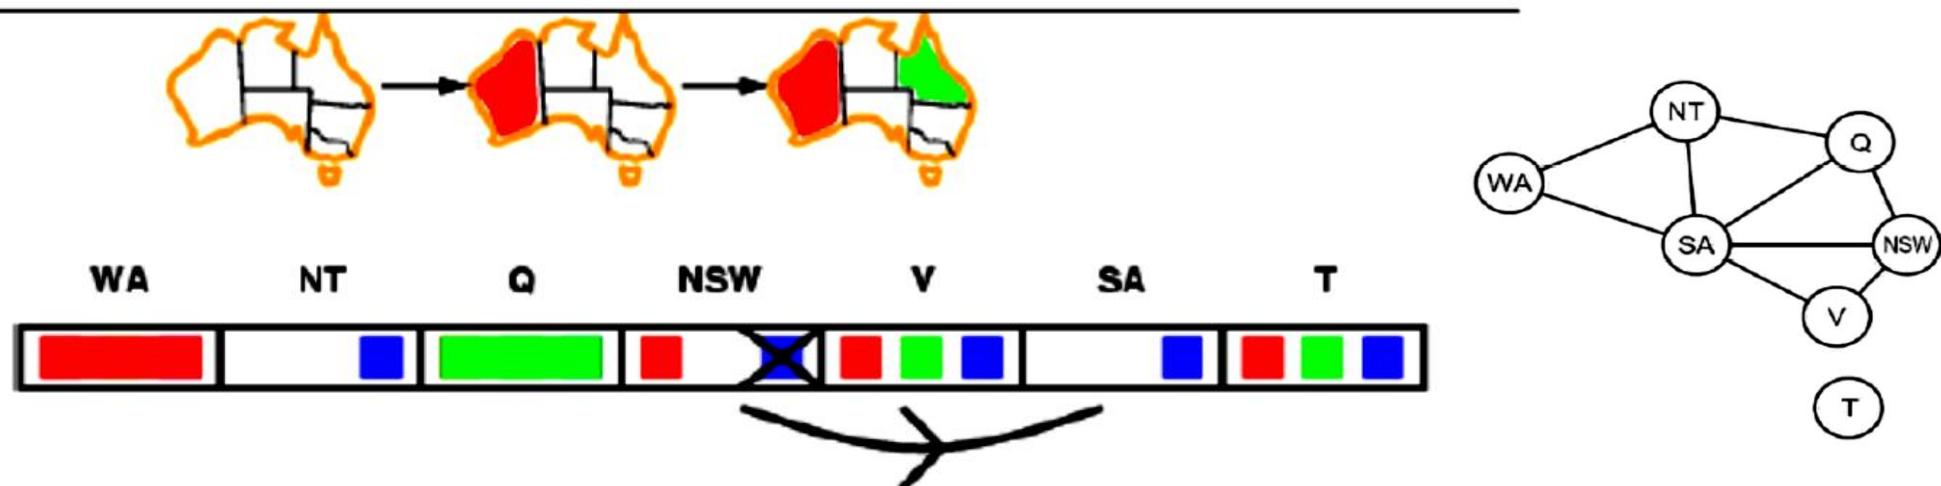
Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



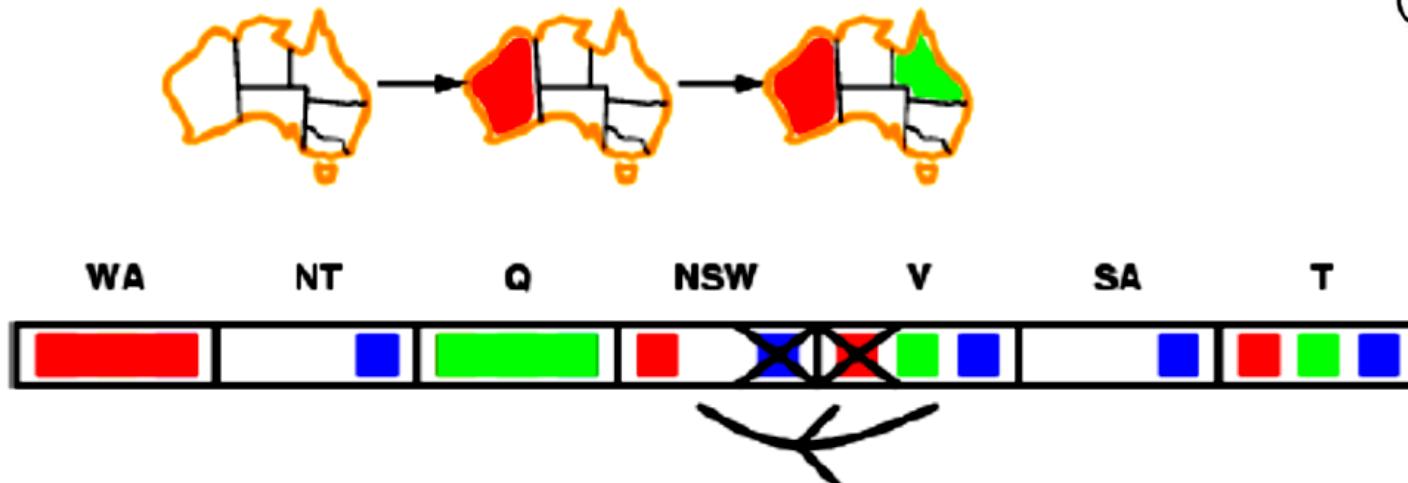
Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



Arc consistency

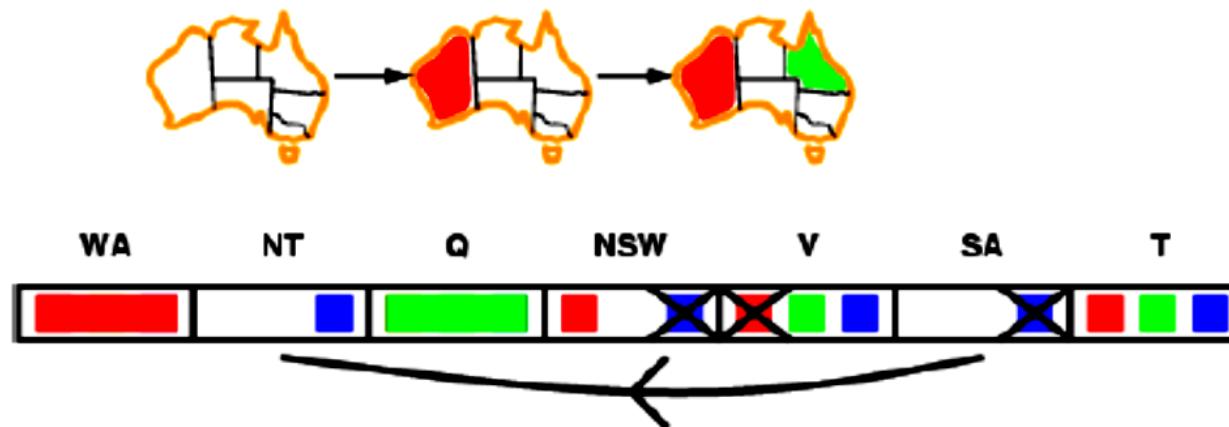
- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



- If X loses a value, neighbors of X need to be rechecked

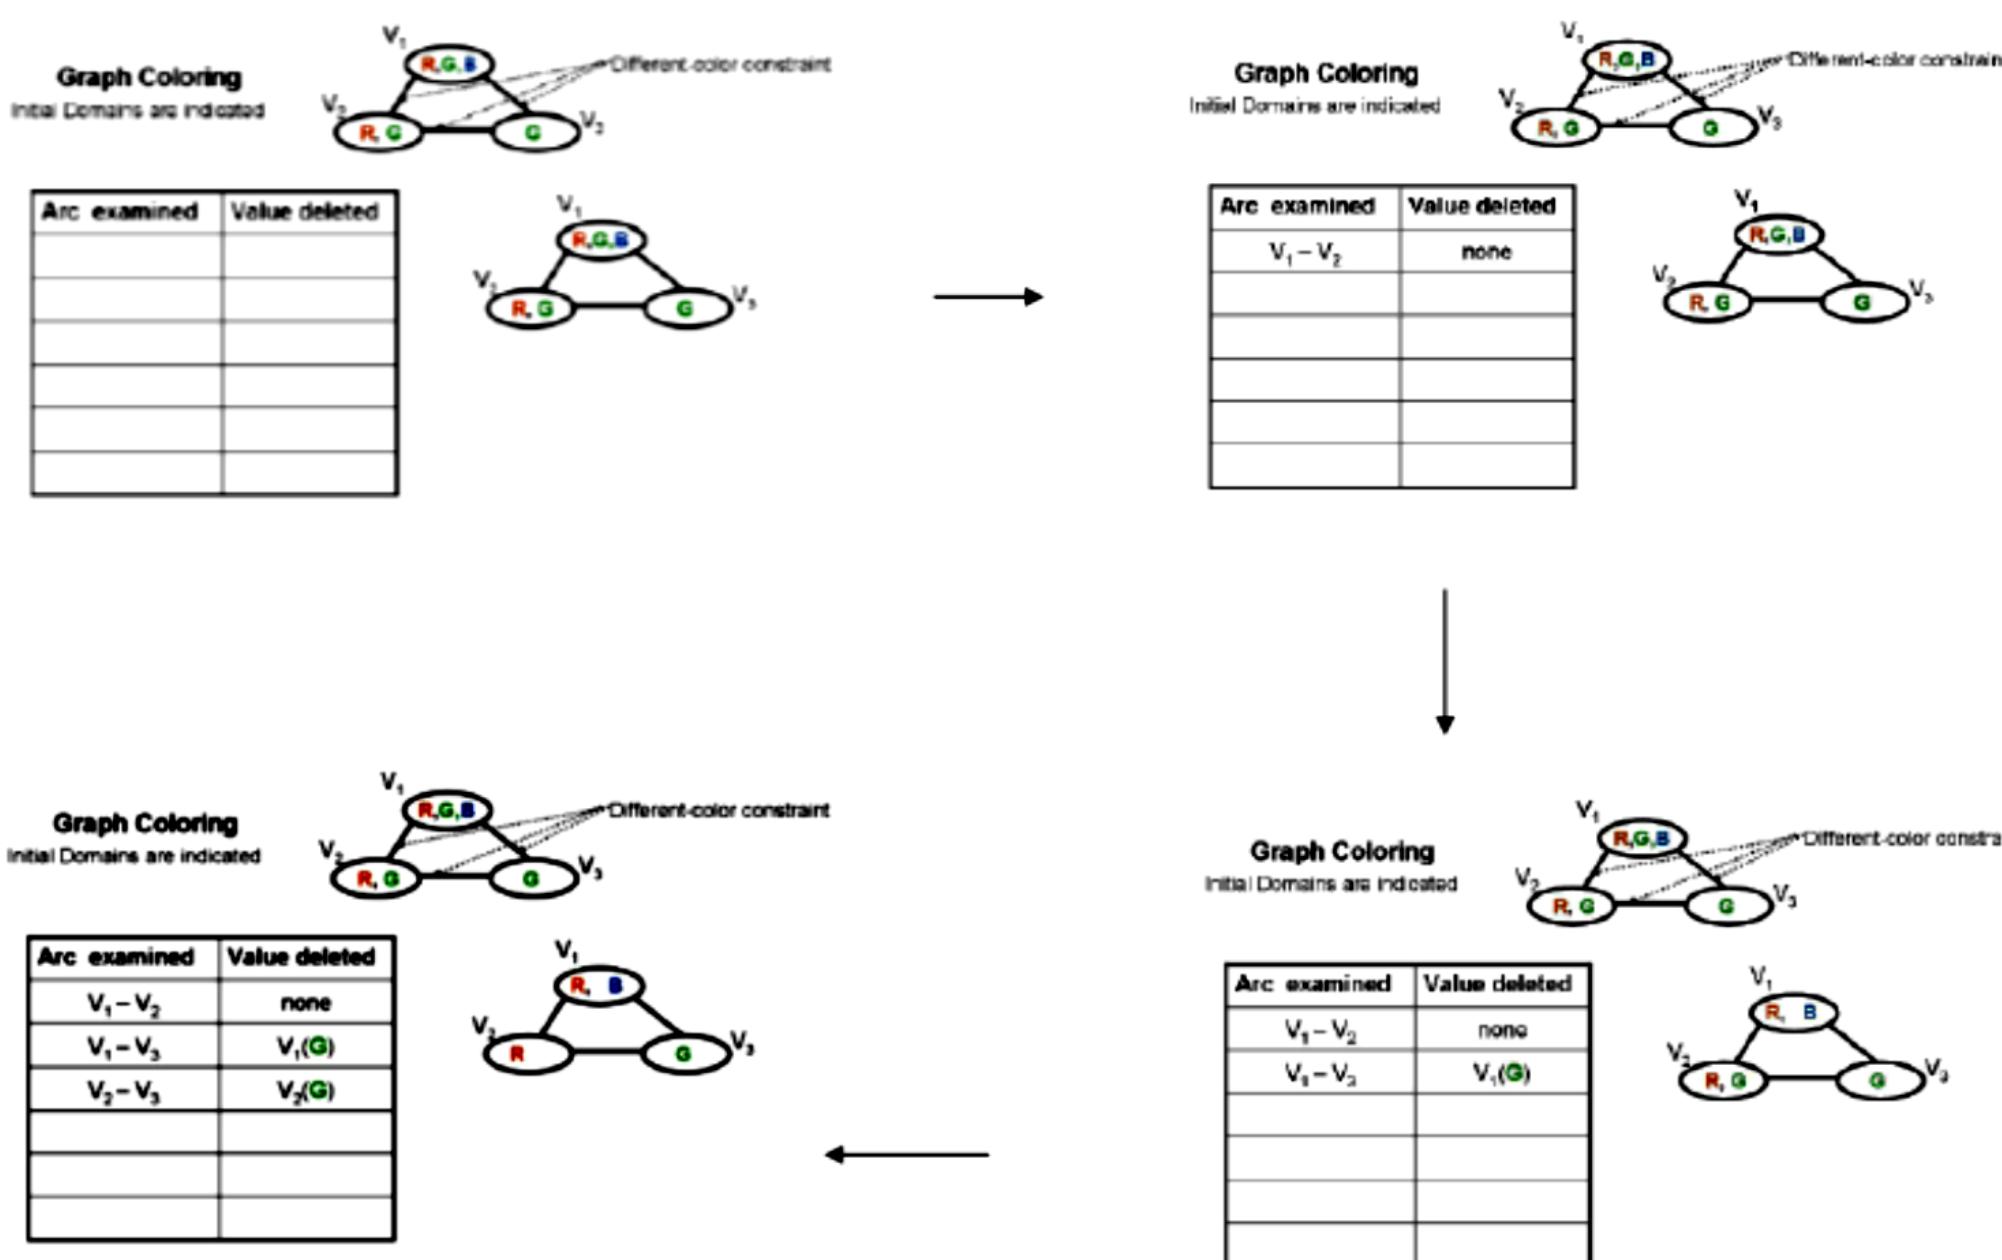
Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \rightarrow Y$ is consistent iff
for **every** value x of X there is **some** allowed y



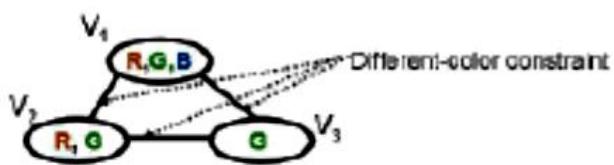
- If X loses a value, neighbors of X need to be rechecked
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment

Constraint Propagation Example 2:

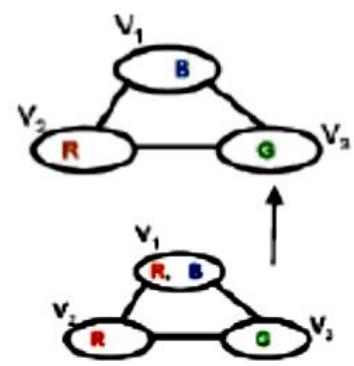


Graph Coloring

Initial Domains are indicated

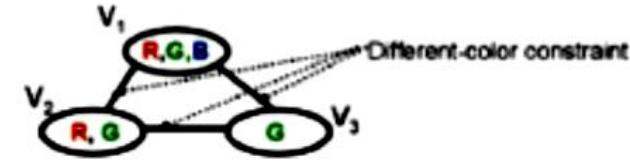


Arc examined	Value deleted
$V_1 - V_2$	none
$V_1 - V_3$	$V_1(G)$
$V_2 - V_3$	$V_2(G)$
$V_1 - V_2$	$V_1(R)$



Graph Coloring

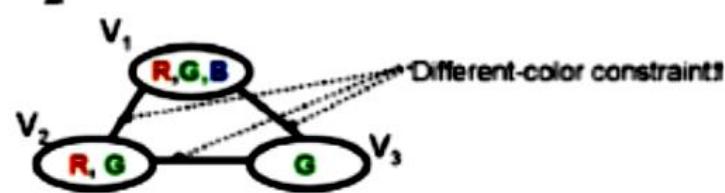
Initial Domains are indicated



Arc examined	Value deleted
$V_1 - V_2$	none
$V_1 - V_3$	$V_1(G)$
$V_2 - V_3$	$V_2(G)$
$V_1 - V_2$	$V_1(R)$
$V_1 - V_3$	none

Graph Coloring

Initial Domains are indicated



Arc examined	Value deleted
$V_1 - V_2$	none
$V_1 - V_3$	$V_1(G)$
$V_2 - V_3$	$V_2(G)$
$V_1 - V_2$	$V_1(R)$
$V_1 - V_3$	none
$V_2 - V_3$	none



Arc consistency algorithm AC-3

```
function AC-3( csp) returns the CSP, possibly with reduced domains
    inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
    local variables: queue, a queue of arcs, initially all the arcs in csp
    while queue is not empty do
         $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
        if RM-INCONSISTENT-VALUES( $X_i, X_j$ ) then
            for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
                add  $(X_k, X_i)$  to queue

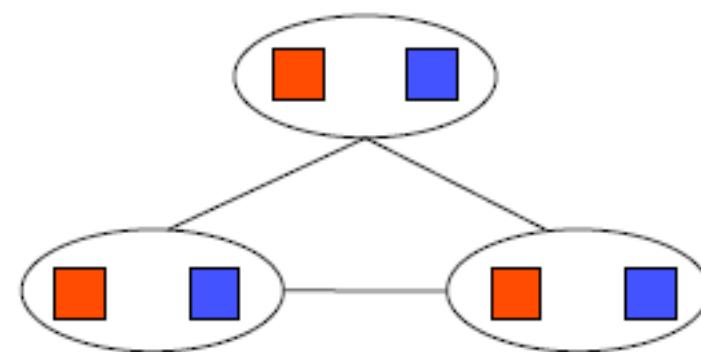
function RM-INCONSISTENT-VALUES(  $X_i, X_j$ ) returns true iff remove a value
    removed  $\leftarrow \text{false}$ 
    for each  $x$  in DOMAIN[ $X_i$ ] do
        if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x,y)$  to satisfy constraint( $X_i, X_j$ )
            then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow \text{true}$ 
    return removed
```

- Time complexity: $O(n^2d^3)$

The story so far:

- CSPs are a special kind of problem:
 - states defined by values of a fixed set of variables
 - goal test defined by constraints on variable values
- Backtracking = depth-first search with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that guarantee later failure
- Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies

Limitations of Arc Consistency



Arc consistent but no sloution

K-Consistency

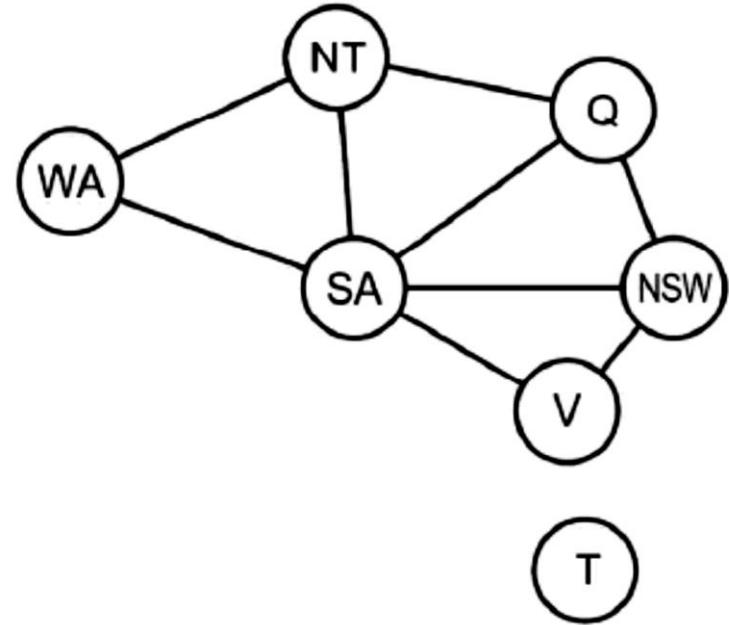
- Increasing degrees of consistency
- 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
- 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
- K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the kth node.
- Higher k more expensive to compute

Trade-offs

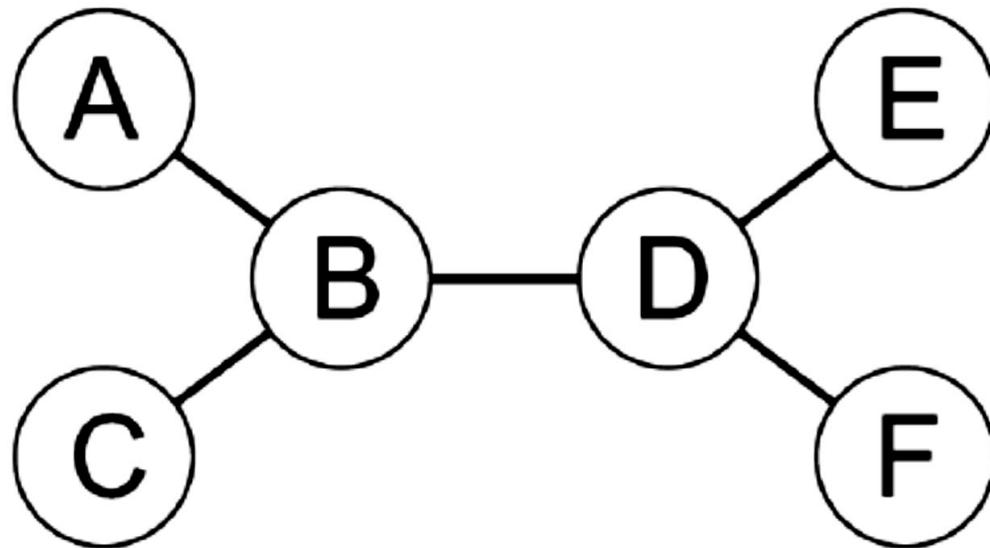
- Running stronger consistency checks...
 - Takes more time
 - But will reduce branching factor and detect more inconsistent partial assignments
 - No “free lunch”
 - In worst case n-consistency takes exponential time
- Generally helpful to enforce 2-Consistency (Arc Consistency)
- Sometimes helpful to enforce 3-Consistency
- Higher levels may take more time to enforce than they save.

Exploiting Graph Structure

- Extreme case: independent subproblems
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as connected components of constraint graph
- Suppose a graph of n variables can be broken into sub-problems of only c variables:
(assume c constant)
 - Worst-case solution cost is $O((n/c)(d^c))$, linear in n
 - E.g., $n = 80$, $d = 2$, $c = 20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $(4)(2^{20}) = 0.4$ seconds at 10 million nodes/sec



Tree-Structured CSPs



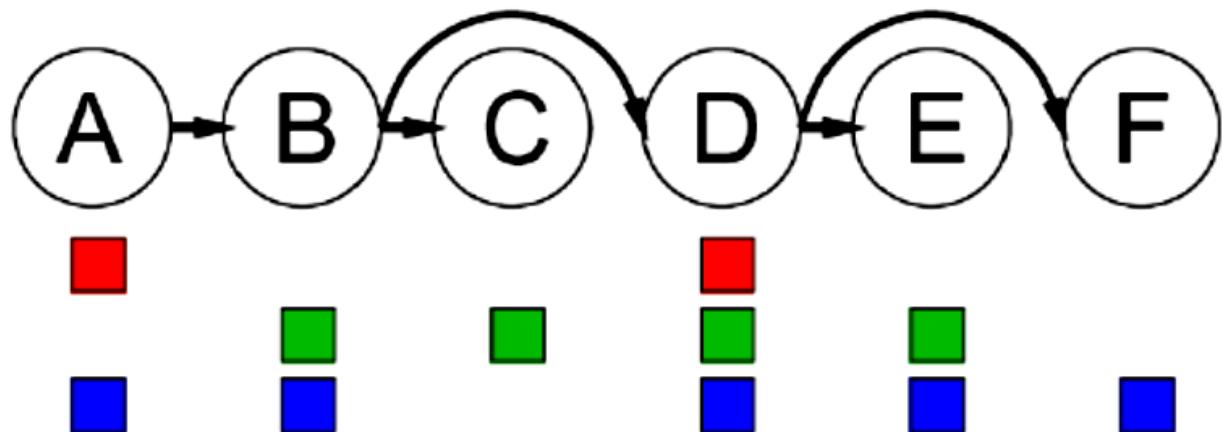
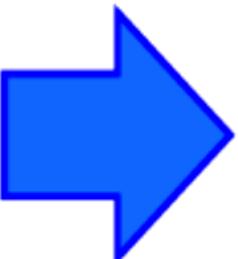
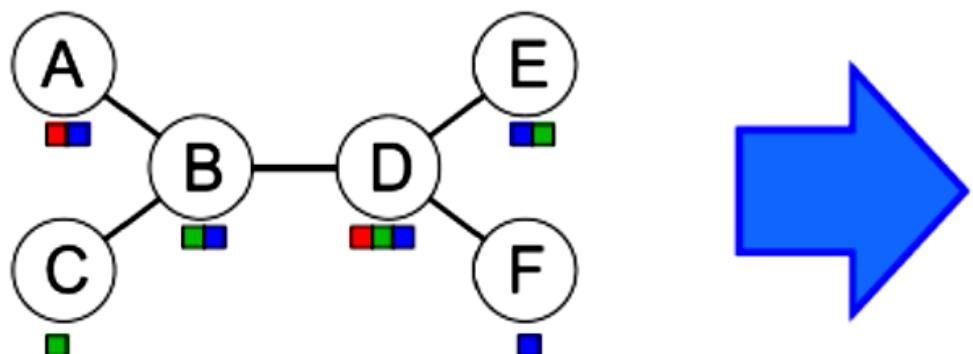
- Theorem: if the constraint graph has no loops, the CSP can be solved in $O(n d^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$

This is linear in n , the number of variables!

Algorithm for Tree-Structured CSPs

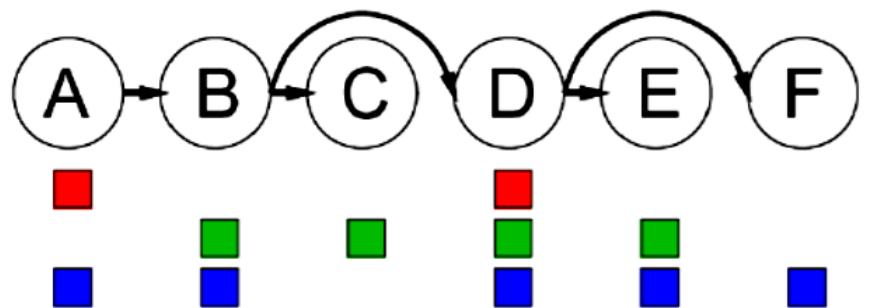
- **Algorithm for tree-structured CSPs:**

- Order: Choose a root variable, order variables so that parents precede children

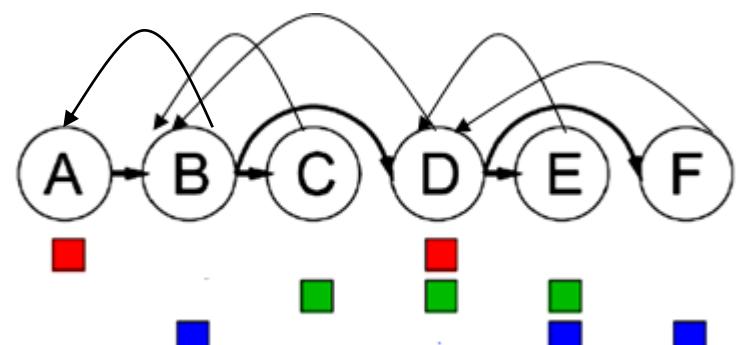
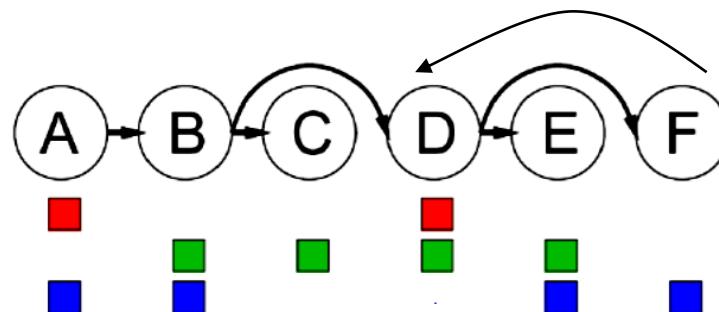


- Backward Pass
 - For j from n down to 2, apply arc consistency to arc $[\text{Parent}(X_j), X_j]$
 - Remove values from $\text{Parent}(X_j)$ if needed
 - Forward Pass
 - For j from 1 to n assign X_j consistently with $\text{Parent}(X_j)$

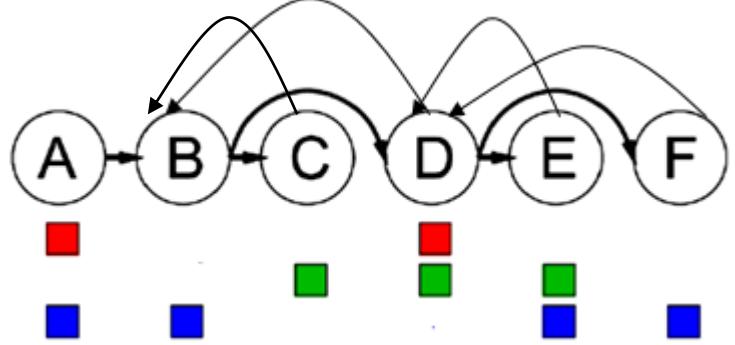
Backward Pass in Action



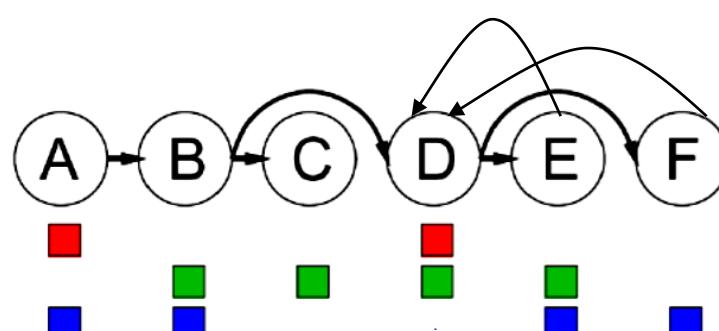
1 →



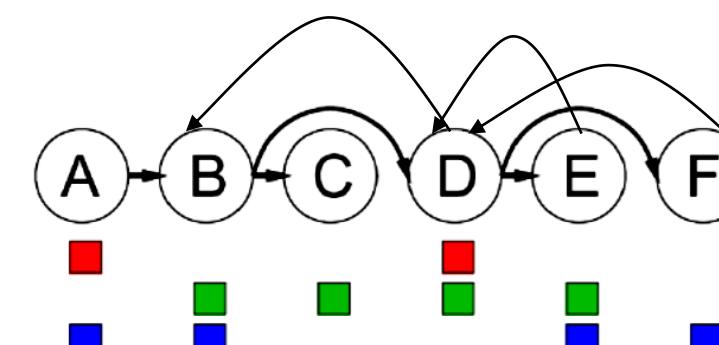
5 ↑



4 ←

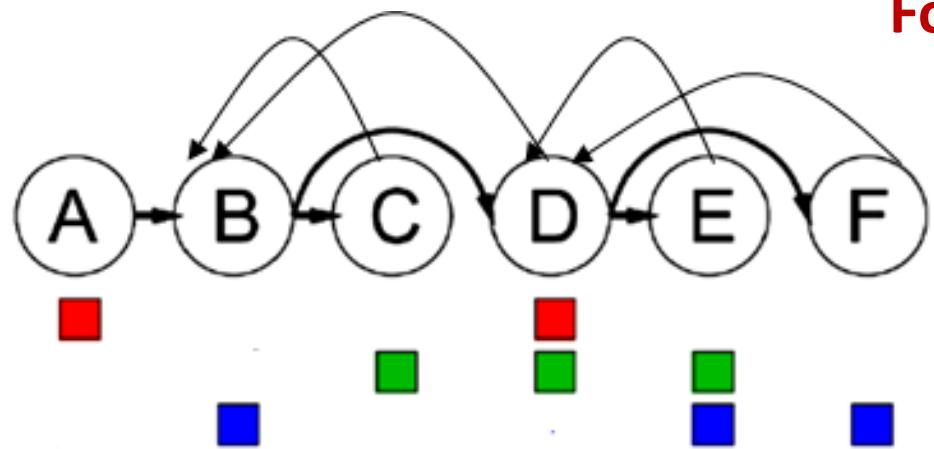


2 ↓

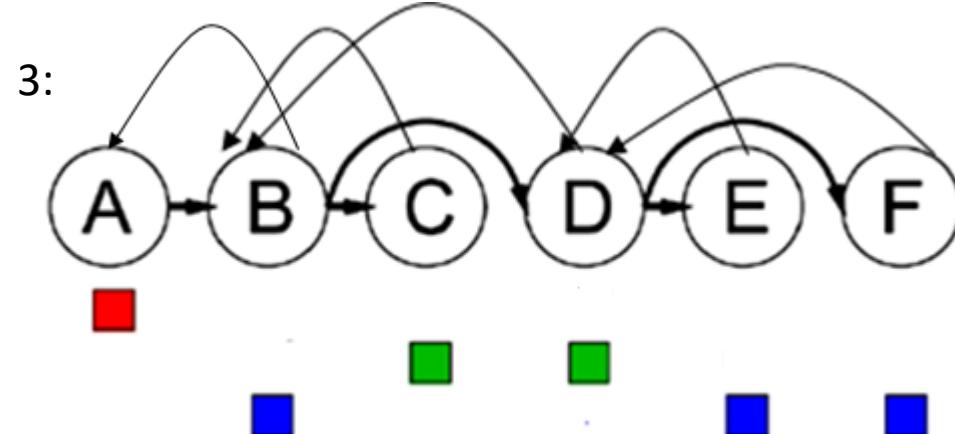
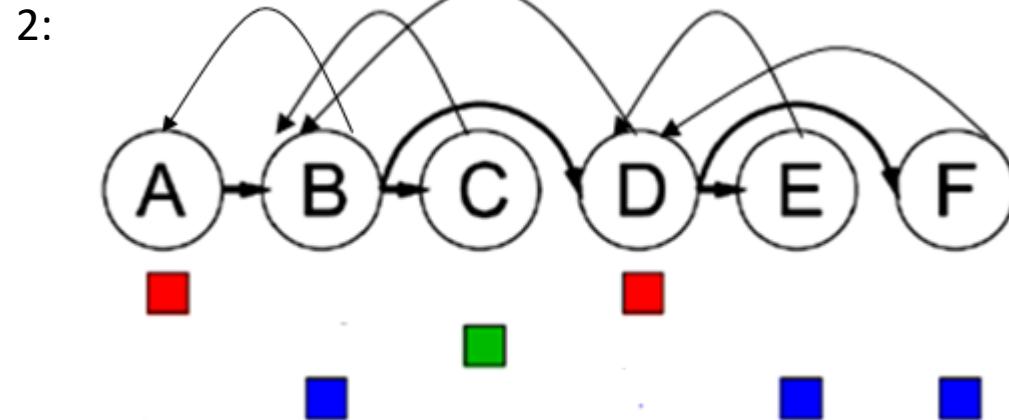
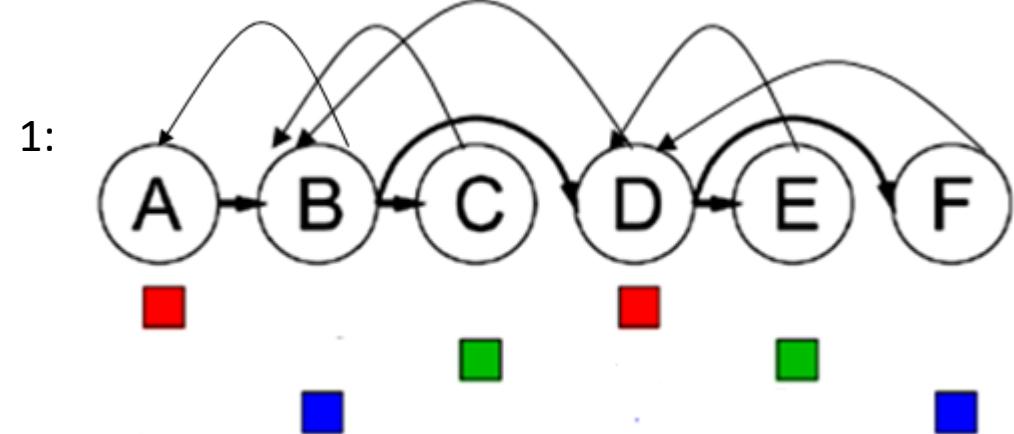


3 ↓

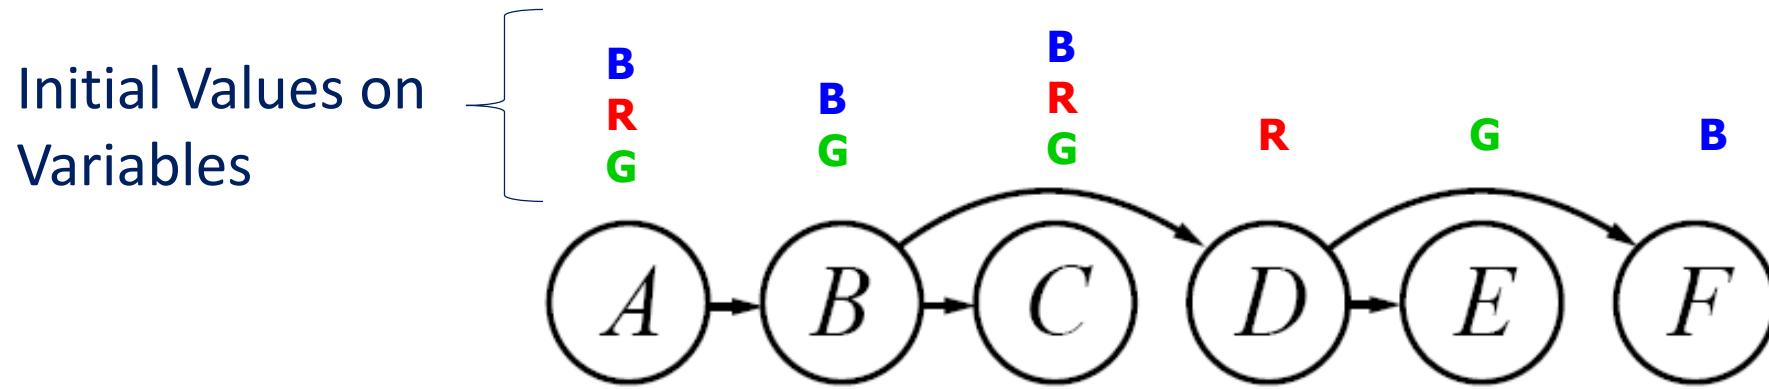
Forward Pass in Action: (3 solutions)



Result of Backward Pass

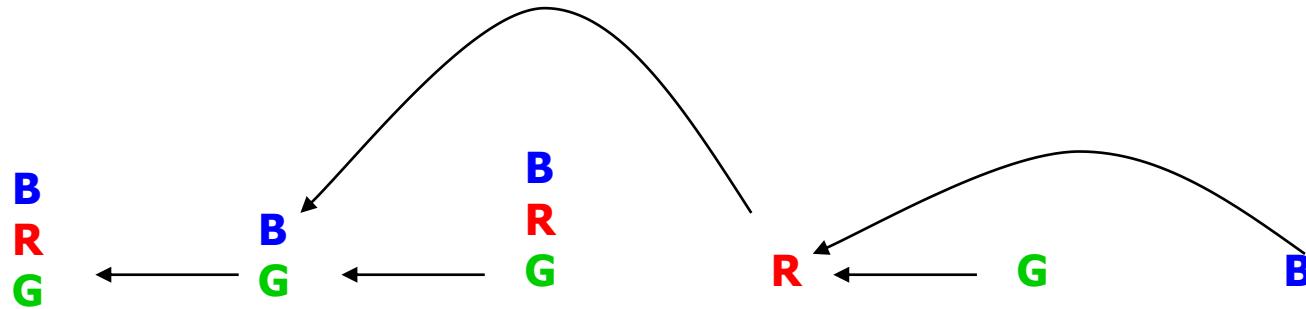


Tree CSP in Action

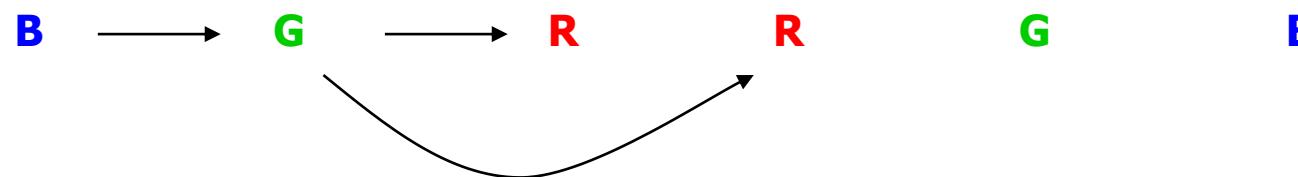


Tree CSP in Action:

Backward Pass
(constraint propagation)

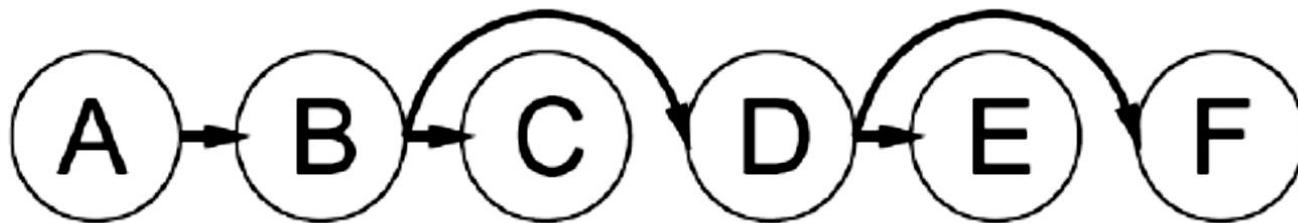


Forward Pass
(assignment)



Why does it work?

- Claim 1: After backward pass, all root-to-leaf arcs are consistent
- Proof: Each $X \rightarrow Y$ was made consistent at one point and Y 's domain could not have been reduced thereafter (because Y 's children were processed before Y)



- Claim 2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
- Proof: Induction on position
- Why doesn't this algorithm work with cycles in the constraint graph?

Tree CSP Complexity

- Backward pass
 - n arc checks
 - Each has complexity d^2 at worst
 - Forward pass
 - n variable assignments, $O(nd)$
- ⇒ Overall complexity is $O(nd^2)$

Algorithm works because if the backward pass succeeds, then every variable by definition has a legal assignment in the forward pass

What about non-tree CSPs?

- General idea is to convert the graph to a tree

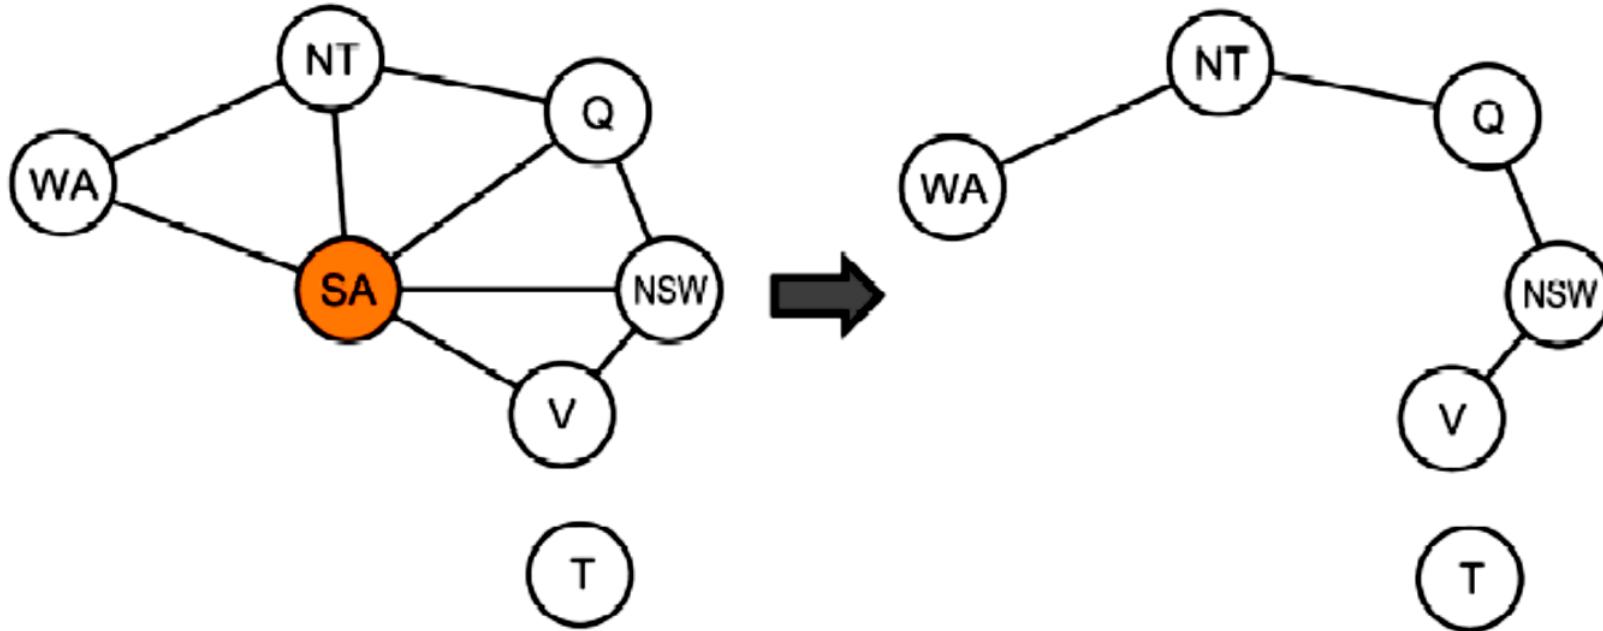
2 general approaches:

- 1) Assign values to specific variables (Cycle Cutset method)
- 1) Construct a tree-decomposition of the graph
 - Connected subproblems (subgraphs) form a tree structure

Cycle-cutset conditioning:

- Choose a subset S of variables from the graph so that graph without S is a tree
 - S = “cycle cutset”
- For each possible consistent assignment for S
 - Remove any inconsistent values from remaining variables that are inconsistent with S
 - Use tree-structured CSP to solve the remaining tree-structure
 - If it has a solution, return it along with S
 - If not, continue to try other assignments for S

Nearly Tree Structured CSPs:



- Conditioning: instantiate a variable, prune its neighbors' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size c gives runtime $O((d^c) (n-c) d^2)$, very fast for small c

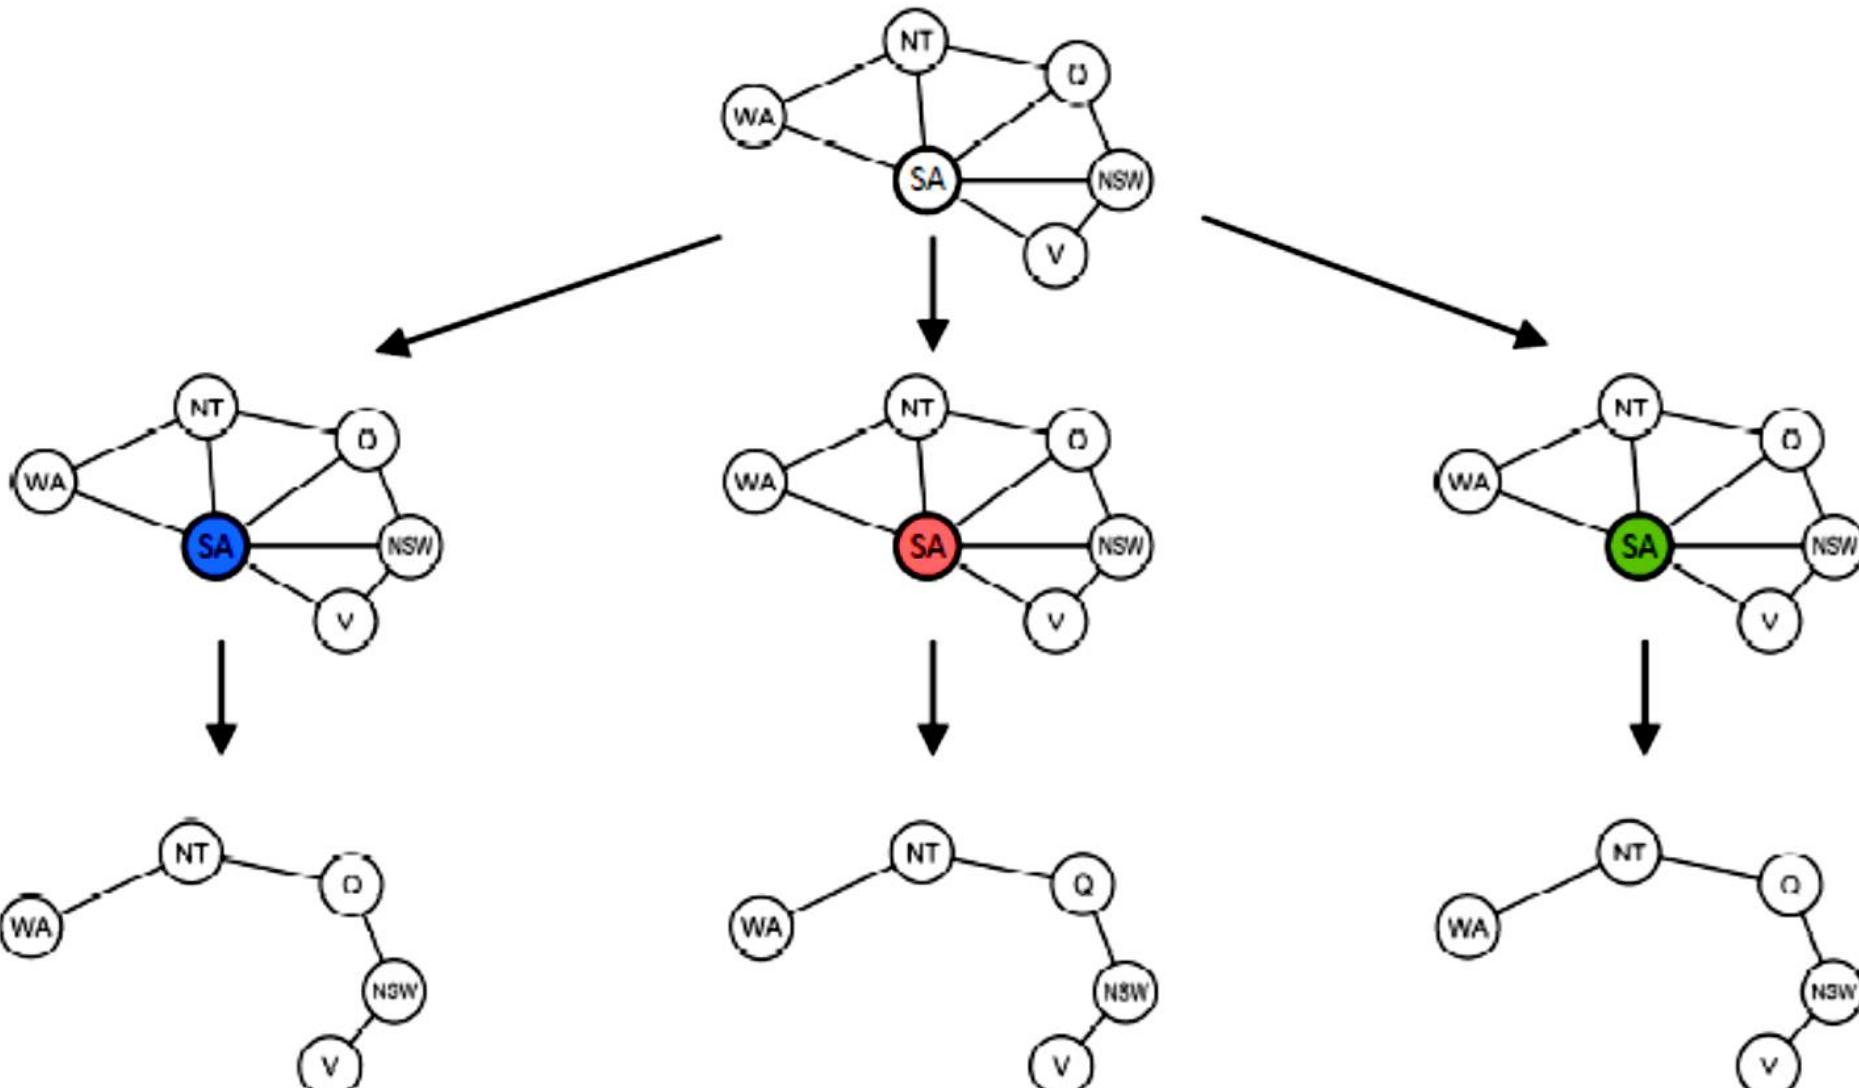
Cutset Conditioning in Action:

Choose a cutset

Instantiate the cutset
(all possible ways)

Compute residual CSP
for each assignment

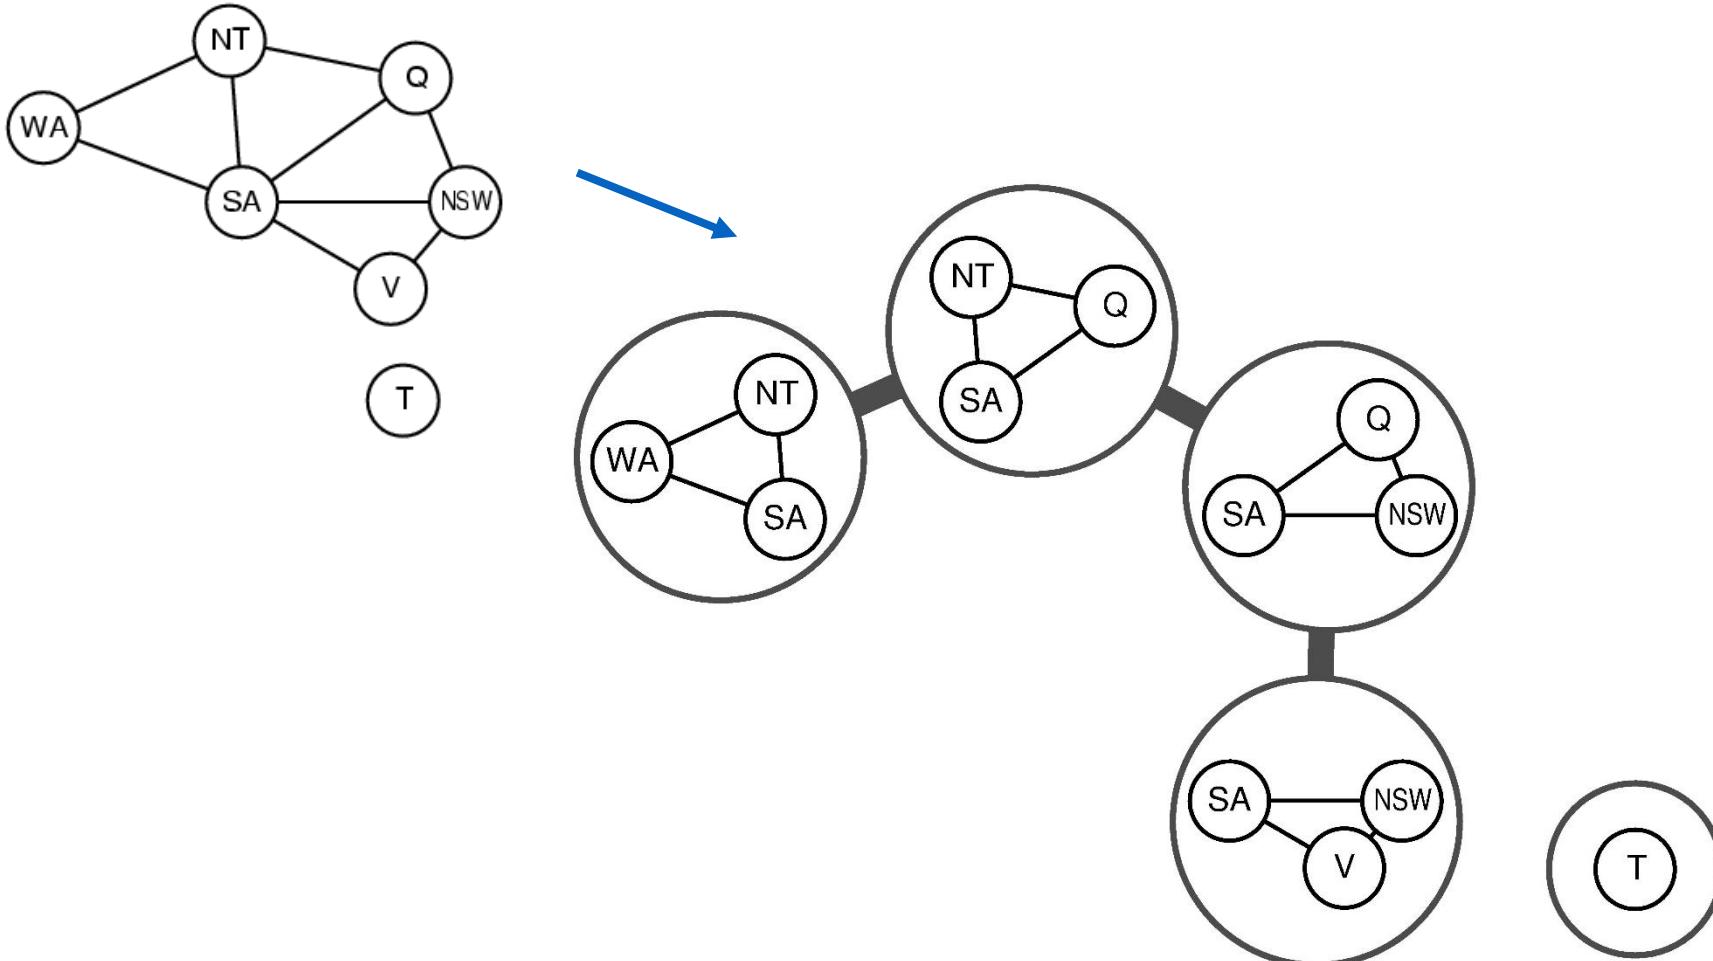
Solve the residual CSPs
(tree structured)



Finding the optimal cutset

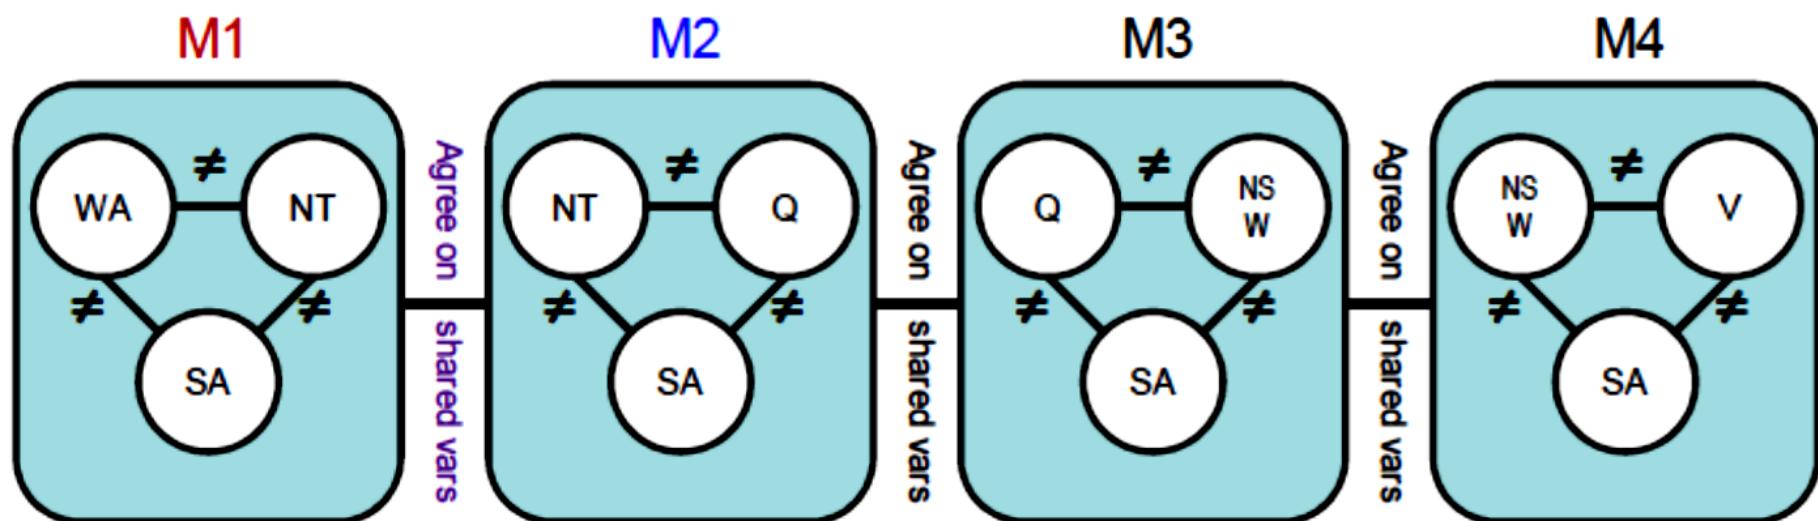
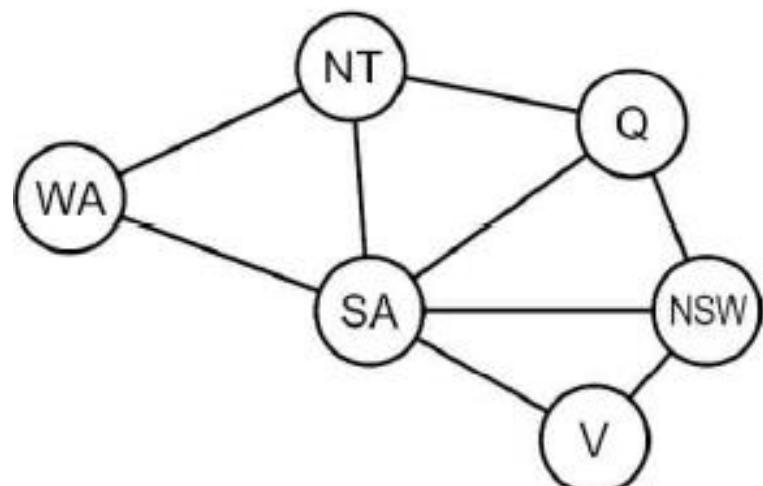
- If c is small, this technique works very well
- However, finding smallest cycle cutset is NP-hard
 - But there are good approximation algorithms

Tree Decomposition:



Note: Just be aware but will not be included in exam

- Idea: create a tree-structured graph of mega-variables
- Each mega-variable encodes part of the original CSP
- Subproblems overlap to ensure consistent solutions



$\{(WA=r, SA=g, NT=b),$
 $(WA=b, SA=r, NT=g),$
 $\dots\}$

$\{(NT=r, SA=g, Q=b),$
 $(NT=b, SA=g, Q=r),$
 $\dots\}$

Agree: $(M1, M2) \in$
 $\{((WA=g, SA=g, NT=g), (NT=g, SA=g, Q=g)), \dots\}$

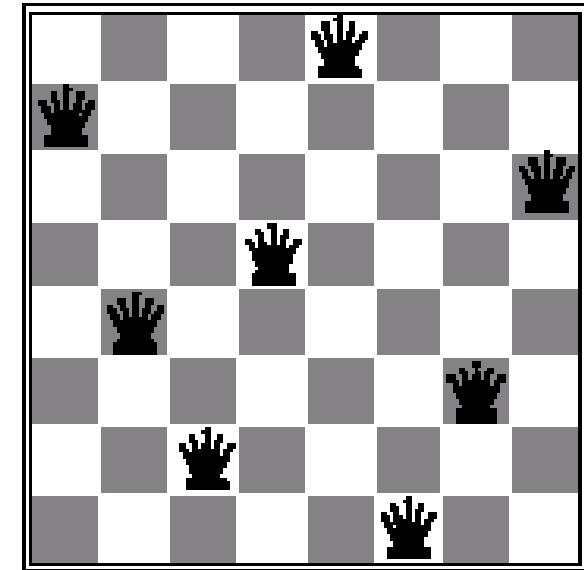
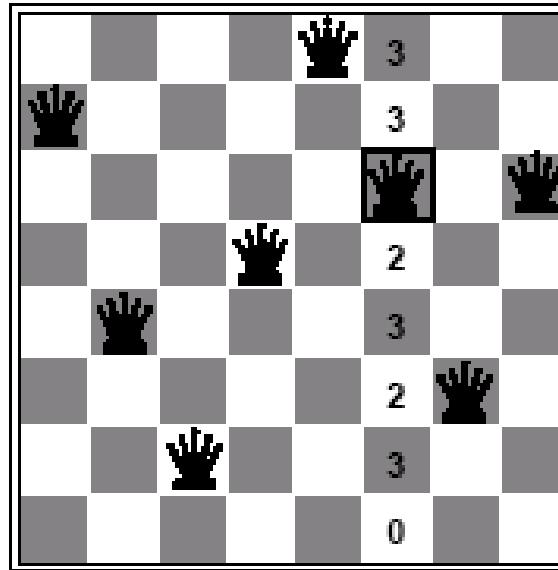
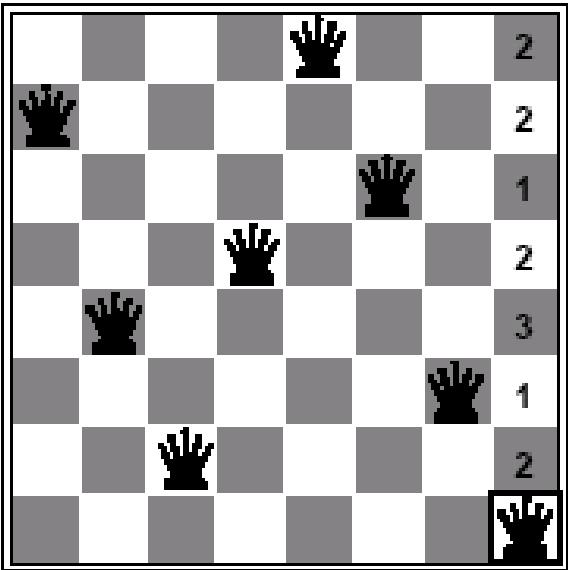
Tree Decomposition Algorithm

- View each subproblem as a “super-variable”
 - Domain = set of solutions for the subproblem
 - Obtained by running a CSP on each subproblem
 - E.g., 6 solutions for 3 fully connected variables in map problem
- Now use the tree CSP algorithm to solve the constraints connecting the subproblems
 - Declare a subproblem a root node, create tree
 - Backward and forward passes
- Example of “divide and conquer” strategy

Local Search a.k.a. Iterative Improvement

- Use complete-state representation
 - Initial state = all variables assigned values
 - Successor states = change 1 (or more) values
- For CSPs
 - allow states with unsatisfied constraints (unlike backtracking)
 - operators reassign variable values
- Variable selection: randomly select any conflicted variable
- Value selection: min-conflicts heuristic
 - Select new value that results in a minimum number of conflicts with the other variables

Min-conflicts Example:



- A two-step solution for an 8-queens problem using min-conflicts heuristic
- At each stage a queen is chosen for reassignment in its column
- The algorithm moves the queen to the min-conflict square breaking ties randomly.

Very powerful in practice: A million-queens in 50 steps average

Time to schedule a week of Hubble Telescope observations cut from 3 weeks to 10 minutes!

Comparison of CSP algorithms on different problems

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA	(> 1,000K)	(> 1,000K)	2K	60	64
n -Queens	(> 40,000K)	13,500K	(> 40,000K)	817K	4K
Zebra	3,859K	1K	35K	0.5K	2K
Random 1	415K	3K	26K	2K	
Random 2	942K	27K	77K	15K	

Median number of consistency checks over 5 runs to solve problem

Parentheses -> no solution found

USA: 4 coloring

n -Queens: n ranges from 2 to 50

CSPs Only Need Binary Constraints!!

Binarization of arbitrary Non-Binary Constraints

$$X \in \{1, 3, 5\}$$

$$X \leq Y$$

$$Y \in \{2, 4, 6\}$$

$$X + Y = Z$$

$$Z \in \{1, 2, 3, 4, 5\}$$

For each non-binary constraint, create an encapsulated variable.

Form the Cartesian Product of the domains of the associated variables,

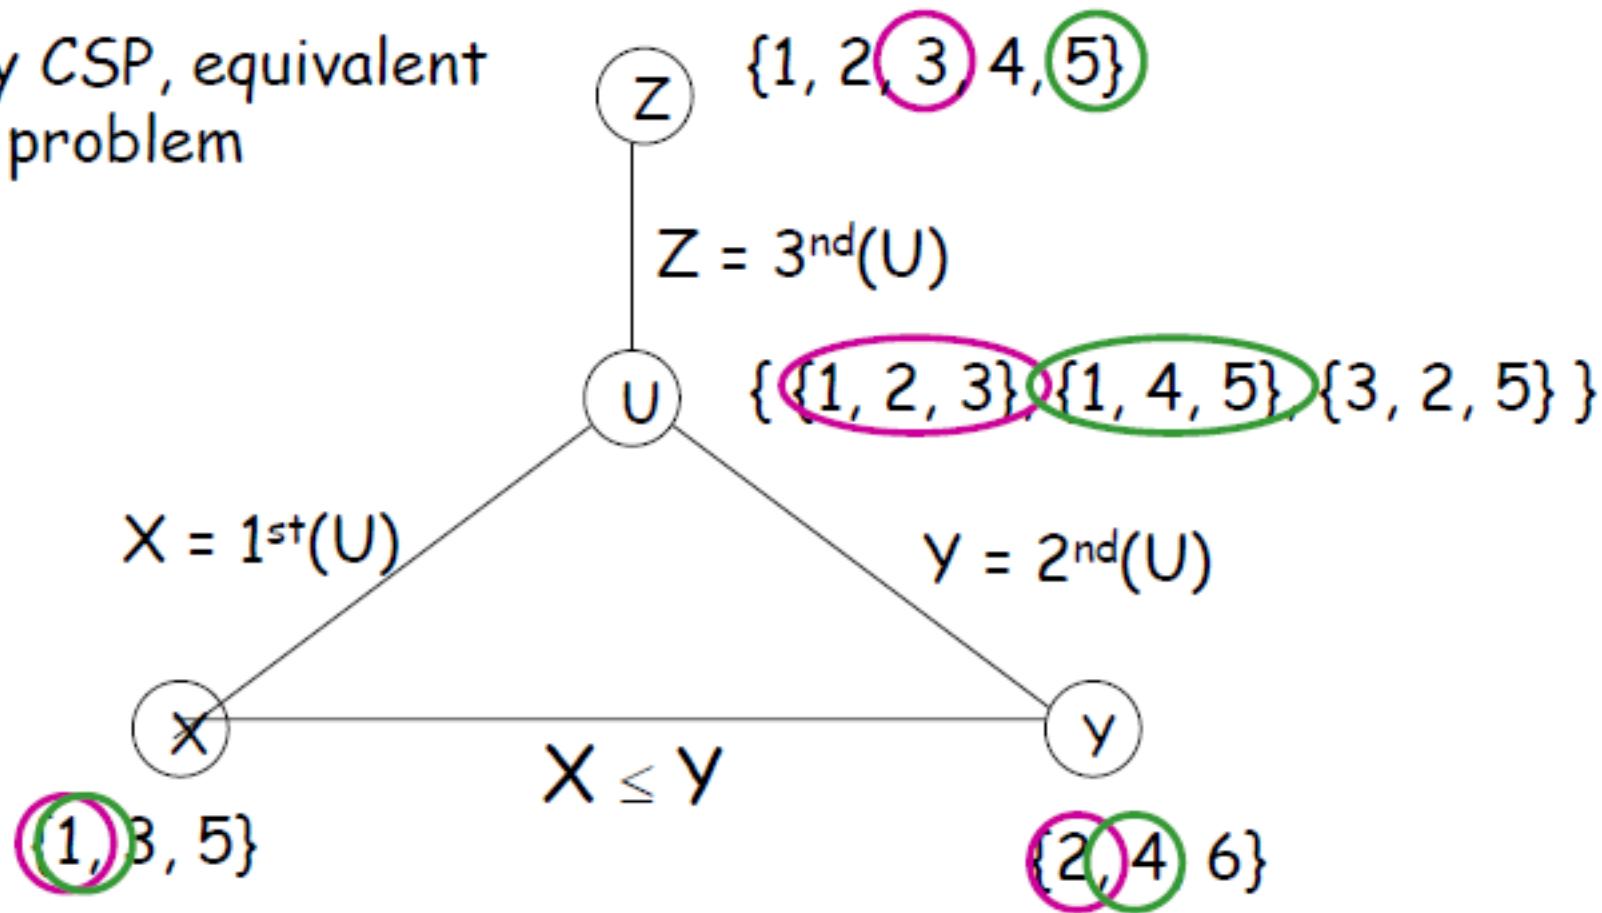
Select those elements which obey the constraint

Cartesian product of X, Y and Z is $\{1, 3, 5\} \times \{2, 4, 6\} \times \{1, 2, 3, 4, 5\}$

The domain of U is a sub-set of the Cartesian product

$$U \in \{\{1, 2, 3\}, \{1, 4, 5\}, \{3, 2, 5\}\}$$

This is a binary CSP, equivalent to the original problem



Preserves the original variables
so solution to original problem can be obtained directly
from the solution to this problem